Thoth

How to find the best application stack

Thoth

How to find the best application stack

How to give the right DNA to our applications

Fridolín Pokorný <fridolin@redhat.com>

Let's develop a machine learning

application

What programming language should I

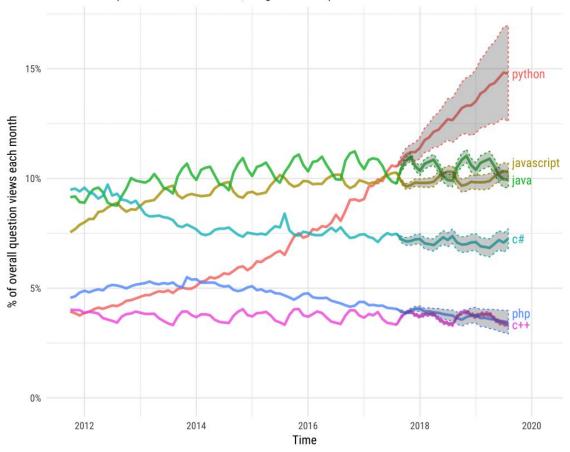
use?

For the sixth year in a row, JavaScript is the most commonly used programming language. Python has risen in the ranks, surpassing C# this year, much like it surpassed PHP last year. Python has a solid claim to being the fastest-growing major programming language.

We see close alignment in the technology choices of professional developers and the developer population overall.

Projections of future traffic for major programming languages

Future traffic is predicted with an STL model, along with an 80% prediction interval.



Source: The Incredible Growth of Python















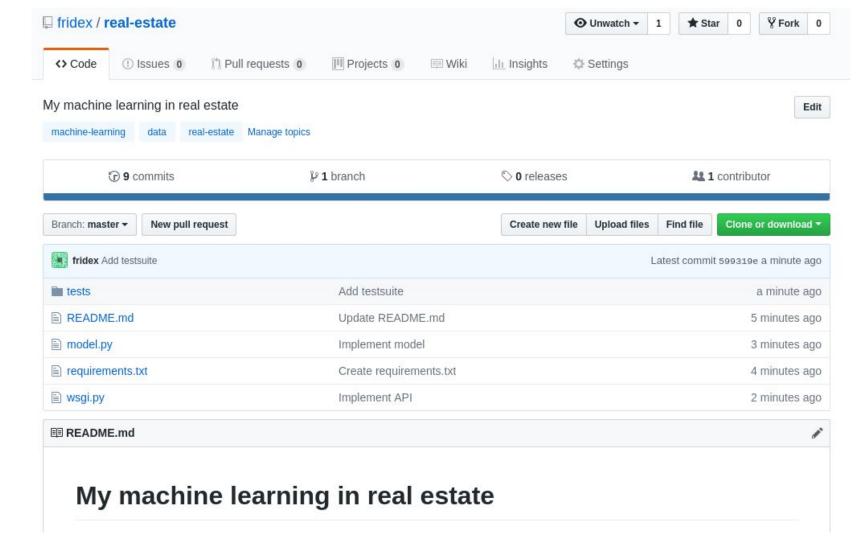


Requirements of my application

requirements.txt file

tensorflow pandas flask gunicorn

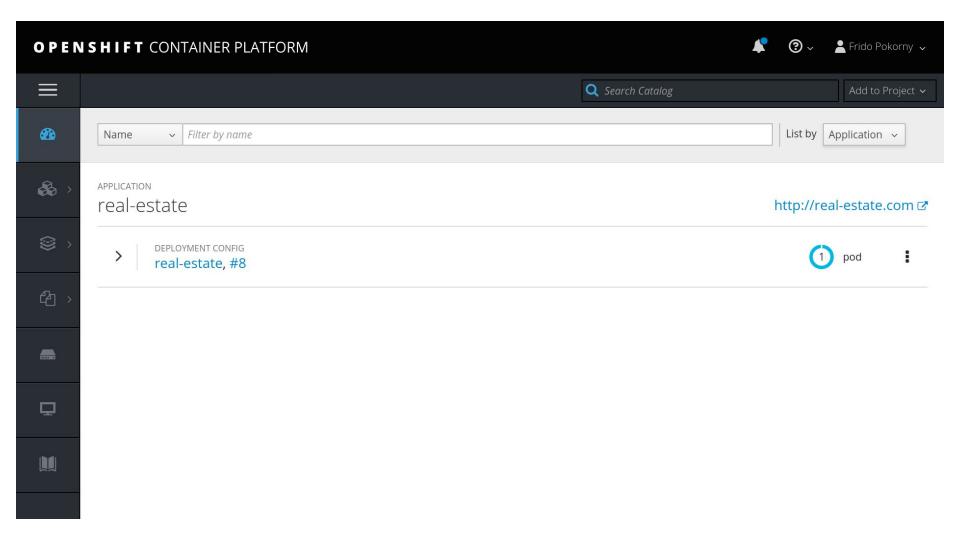




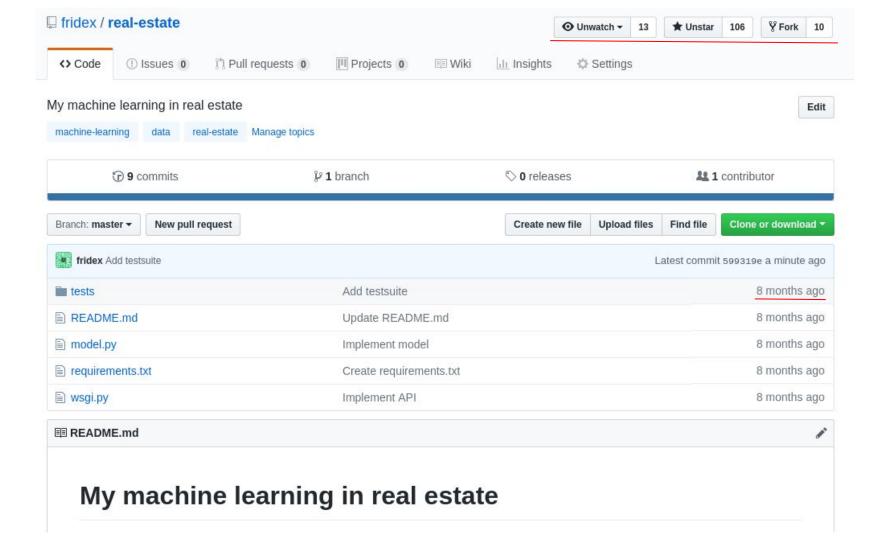


source to image







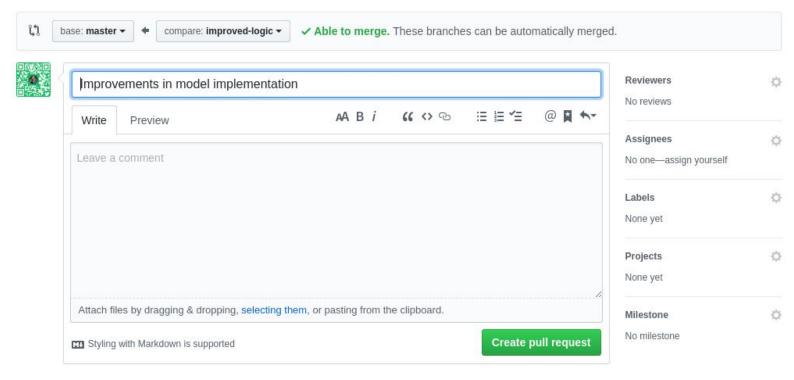






Open a pull request

Create a new pull request by comparing changes across two branches. If you need to, you can also compare across forks.





Jenkins



Status

Changes

Console Output

View as plain text

View Build Information

it Build Data

Replay

- " " - - ·

Pipeline Steps

Previous Build





Build Artifacts

test.log

582.17 KB 📖 view



Started by timer



Revision: 73c310b7035480c2f146f00358bb85f0b7510c72

· refs/remotes/origin/master

Let's debug application



Behavior change in library

Previous behavior

In [1]: import pandas as pd

In [2]: pd.__version__
Out[2]: '0.21.0'

In [1]: import pandas as pd

```
In [1]: import pandas as pd
In [2]: pd. version
Out[2]: '0.21.0'
In [3]: df = pd.DataFrame([[10, None, 30], []])
        df
Out[3]:
             0 1 2
        0 10.0 None 30.0
        1 NaN None NaN
In [4]: df = df.sum()
        df
Out[4]: 0
            10.0
             NaN
             30.0
        dtype: float64
```

```
In [1]: import pandas as pd
In [2]: pd. version
Out[2]: '0.21.0'
In [3]: df = pd.DataFrame([[10, None, 30], []])
        df
Out[3]:
             0 1 2
        0 10.0 None 30.0
        1 NaN None NaN
In [4]: df = df.sum()
        df
Out[4]: 0
            10.0
             NaN
             30.0
        dtype: float64
In [5]: df.mean()
Out[5]: 20.0
```

Behavior now

```
In [1]: import pandas as pd

In [2]: pd.__version__
Out[2]: '0.22.0'
```

```
In [1]: import pandas as pd

In [2]: pd.__version__
Out[2]: '0.22.0'

In [3]: df = pd.DataFrame([[10, None, 30], []])
df
```

```
In [1]: import pandas as pd
In [2]: pd.__version__
Out[2]: '0.22.0'
In [3]: df = pd.DataFrame([[10, None, 30], []])
        df
Out[3]:
             0 1 2
        0 10.0 None 30.0
        1 NaN None NaN
```

```
In [1]: import pandas as pd
In [2]: pd. version
Out[2]: '0.22.0'
In [3]: df = pd.DataFrame([[10, None, 30], []])
        df
Out[3]:
             0 1 2
        0 10.0 None 30.0
        1 NaN None NaN
In [4]: df = df.sum()
        df
Out[4]: 0
            10.0
             0.0
            30.0
        dtype: float64
```

```
In [2]: pd. version
Out[2]: '0.22.0'
In [3]: df = pd.DataFrame([[10, None, 30], []])
        df
Out[3]:
             0 1 2
        0 10.0 None 30.0
        1 NaN None NaN
In [4]: df = df.sum()
        df
Out[4]: 0
            10.0
            0.0
            30.0
        dtype: float64
In [5]: df.mean()
Out[5]: 13.3333333333333334
```

In [1]: import pandas as pd

Change in behavior across versions

What versions of libraries do I use to develop my application?

What versions of libraries is my application compatible with?

- What if libraries I use change over time?
 - They change over time as can be seen!



should i pin down my python dependencies versions



Google Search

I'm Feeling Lucky



should I pin down my python dependencies versions





All

Videos

Images

News

Shopping

More

Settings

Tools

About 458,000 results (0.57 seconds)

Should I pin my Python dependencies versions? - Stack Overflow

https://stackoverflow.com/questions/.../should-i-pin-my-python-dependencies-version... ▼ 3 answers

Feb 14, 2015 - You should always pin your dependencies as it increases the possibility of ... is to version your dependencies, as the relationship of the PEP on ...

python - Why should version numbers not be pinned in a pipfile ... Jun 13, 2018

python - Is it possible to lock versions of packages in Anaconda ... Feb 10, 2018

packaging - How to pin transitive dependencies with setup.py in ... Jun 9, 2017

python - Installing specific package versions with pip Mar 7, 2011

More results from stackoverflow.com

Pin Your Packages » nvie.com

https://nvie.com/posts/pin-your-packages/ >

Sep 26, 2012 - Update: A newer blog post about the future of pip-tools is available too: ... a) what's currently installed, and b) what's the current version on PyPI. Eventually, all of your environments, and those of your team members, will run ...

To pin or not to pin dependencies: reproducible vs. reusable software

blog.chrisgorgolewski.org/2017/12/to-pin-or-not-to-pin-dependencies.html ▼
Dec 3, 2017 - On the other side when you opt not to pin dependencies (for ... the package dependencies or create a new Python environment just to use your package. ... you should only specify minimally required dependency versions and ...

Should I pin my Python dependencies versions?



I am about to release a Python library I've been working on the past few weeks. I've read a lot about Python dependencies but something is not quite clear yet:

12

Some people pretend you should never pin your dependencies versions as it would prevent the users of your library from upgrading those dependencies.



Some other claim that you should always pin your dependencies versions as it is the only way of guaranteeing that your release works the way it did when you developped it and to prevent that a breaking change in a dependency wreaks havoc in your library.

I'm somehow went for an hybrid solution, where I assumed my dependencies used semantic versioning and pinned only the major version number (say somelib >= 2.3.0, < 3) except when the major version number is 0 (semantic versioning dictates that such versions are to be considered volatile and may break the API even if only the patch number is bumped).

As of now, I'm not sure which way is the best. Is there an official guideline (even a PEP perhaps?) that dictates the best practice regarding Python dependencies and how to specify them?





3 Answers votes



You should always pin your dependencies as it increases the possibility of safe, repeatable builds, even as time passes. The pinned versions are your declaration as a package maintainer that you've verified that your code works in a given environment. This has a nice side effect of preserving your sanity as you won't be inundated with bug reports in which you have to play inspector into every package codependency and system detail.



Users can always choose to ignore the pinned dependency-versions and do so at their own risk. However, as you release new versions of your library, you should update your dependency versions to take in improvements and bug fixes.

The section of PEP 426 about Semantic dependencies (Metadata for Python Software Packages) states:



You should always pin your dependencies as it increases the possibility of safe, repeatable builds, even as time passes. The pinned versions are your declaration as a package maintainer that you've verified that your code works in a given environment. This has a nice side effect of preserving your sanity as you won't be inundated with bug reports in which you have to play inspector into every package codependency and system detail.



Users can always choose to ignore the pinned dependency-versions and do so at their own risk. However, as you release new versions of your library, you should update your dependency versions to take in improvements and bug fixes.

The section of PEP 426 about Semantic dependencies (Metadata for Python Software Packages) states:

"Dependency management is heavily dependent on the version identification and specification scheme defined in PEP 440 (PEP 440 - Version Identification and Dependency Specification)."

From this, I infer that the authoritative "best practice" is to version your dependencies, as the relationship of the PEP on packaging is stated to be "heavily dependent" on the versioning details outlined by the related PEP.

share improve this answer



Requirements of my application

requirements.txt file

tensorflow pandas flask gunicorn

Requirements of my application

requirements.txt file

tensorflow pandas flask gunicorn

My requirements for the application

• requirements.in file

tensorflow pandas flask gunicorn

My requirements for the application

requirements.in file

tensorflow pandas flask gunicorn

Requirements of the application

• requirements.txt file ?

Requirements of application

requirements.txt file

```
abs1-py==0.3.0
                          # via tensorflow
astor==0.7.1
                         # via tensorflow
click==6.7
                          # via flask
flask==0.6.0
qast == 0.2.0
                        # via tensorflow
                          # via tensorflow
grpcio==1.14.1
qunicorn==19.9.0
itsdangerous==0.24
                          # via flask
jinja2==2.10
                         # via flask
markdown==2.6.11
                          # via tensorboard
markupsafe==1.0
                          # via jinja2
numpy = 1.14.5
                          # via pandas, tensorboard, tensorflow
pandas==0.23.4
protobuf==3.6.0
                          # via tensorboard, tensorflow
python-dateutil==2.7.3
                         # via pandas
pvtz==2018.5
                          # via pandas
six = 1.11.0
                          # via absl-py, grpcio, protobuf, python-dateutil, tensorboard, tensorflow
tensorboard==1.10.0
                          # via tensorflow
tensorflow==1.10.0
termcolor==1.1.0
                         # via tensorflow
werkzeug==0.14.1
                        # via flask, tensorboard
                          # via tensorboard, tensorflow
wheel==0.31.1
```

Requirements of application

Fully pinned down software stack

• Transitive dependencies

• Resolution is no more time dependent

Recommended - use pipenv

\$ git commit -m "Pin down dependencies"
\$ git push

▲ We found a potential security vulnerability in one of your dependencies.

\$ git commit -m "Excluded vulnerable version"
\$ git push



Jenkins



Status

Changes

Console Output

View as plain text

View Build Information

it Build Data

Replay

Pipeline Steps

Previous Build



Build #142 (Aug 13, 2018 6:00:00 PM)



Build Artifacts

test.log

582.17 KB 🚞 view



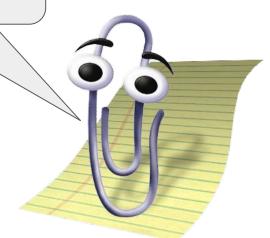
Started by timer



Revision: 73c310b7035480c2f146f00358bb85f0b7510c72

· refs/remotes/origin/master

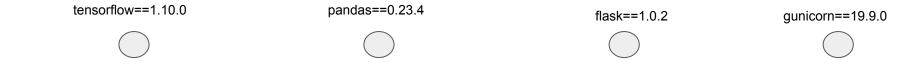
Hey, do you know the fact you are definitely using wrong software?

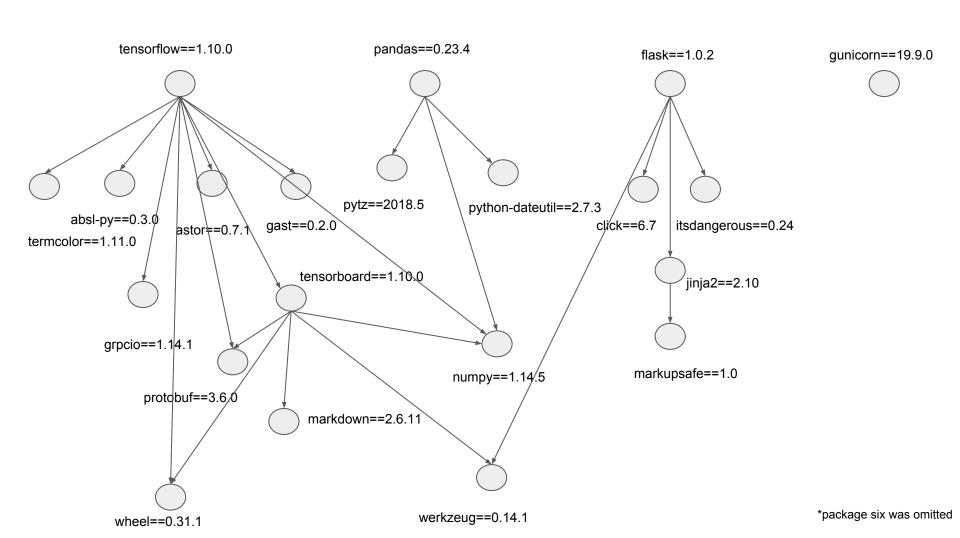


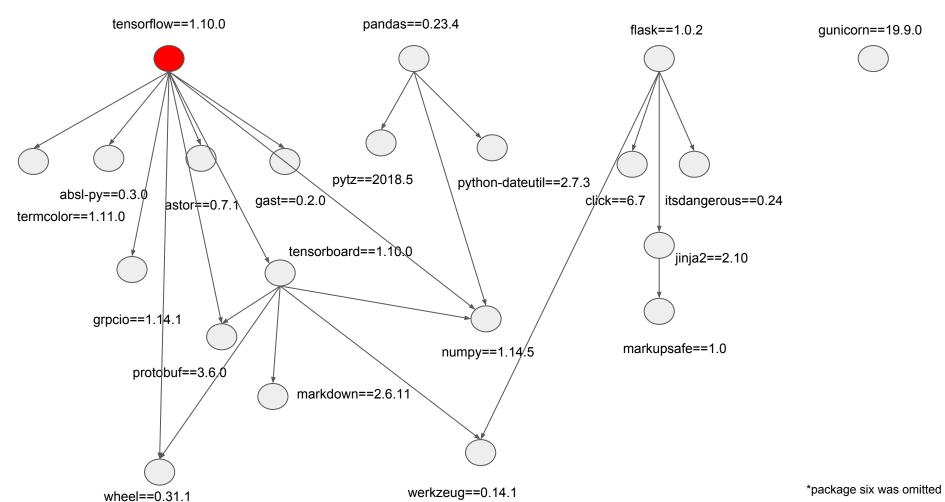
```
tensorflow
pandas
flask
gunicorn
```

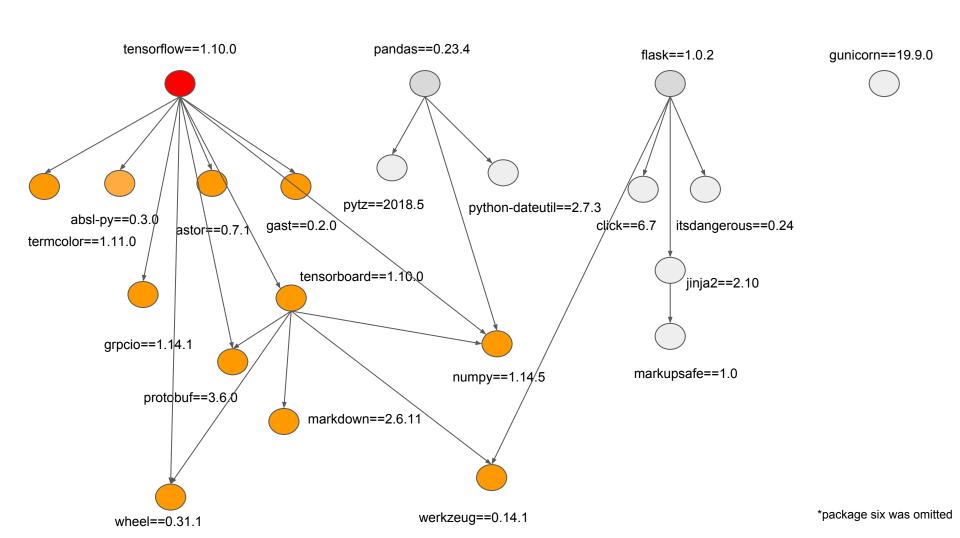
```
absl-py==0.3.0
astor==0.7.1
click==6.7
flask==0.6.0
gast==0.2.0
grpcio==1.14.1
gunicorn==19.9.0
itsdangerous==0.24
jinja2==2.10
markdown==2.6.11
markupsafe==1.0
numpy = 1.14.5
pandas==0.23.4
protobuf==3.6.0
python-dateutil==2.7.3
pytz==2018.5
six = 1.11.0
tensorboard==1.10.0
tensorflow==1.10.0
termcolor==1.1.0
```

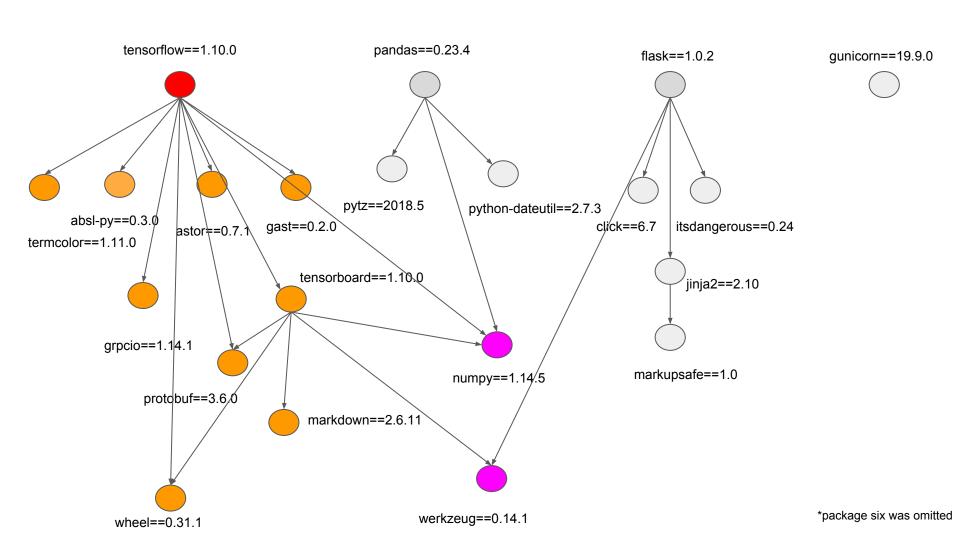
werkzeug==0.14.1 wheel==0.31.1

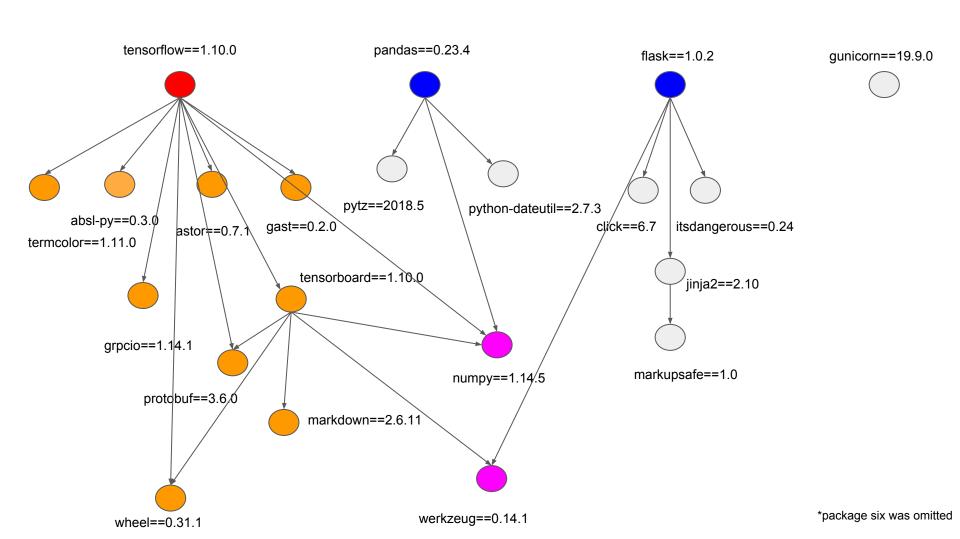


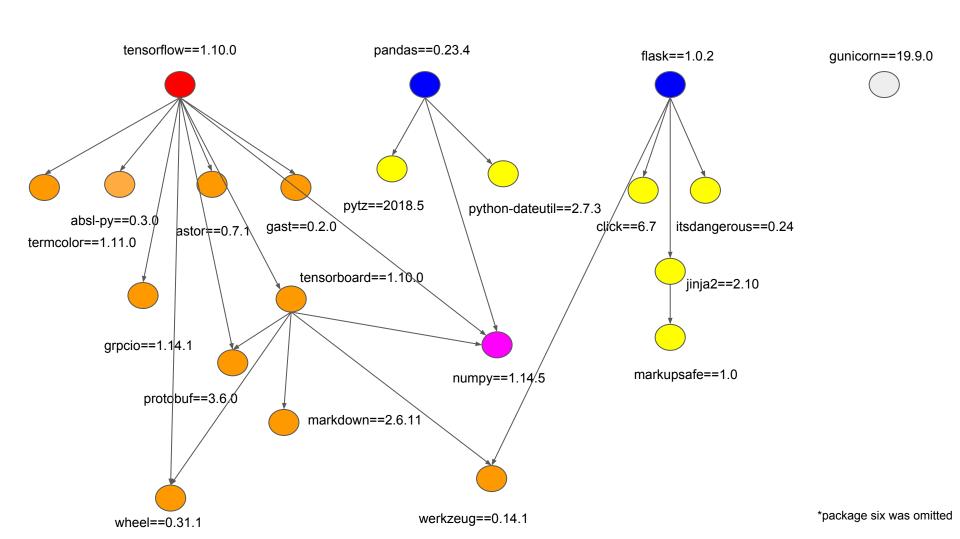


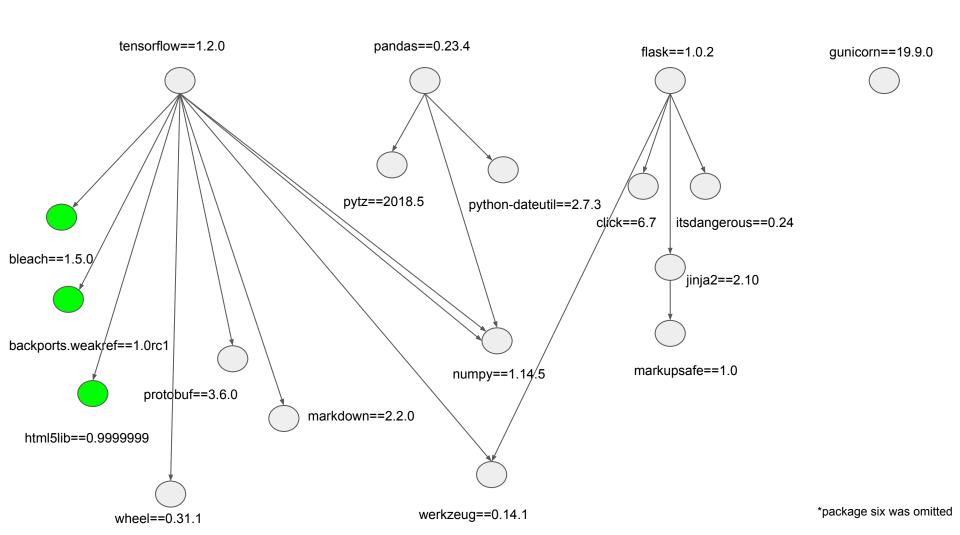


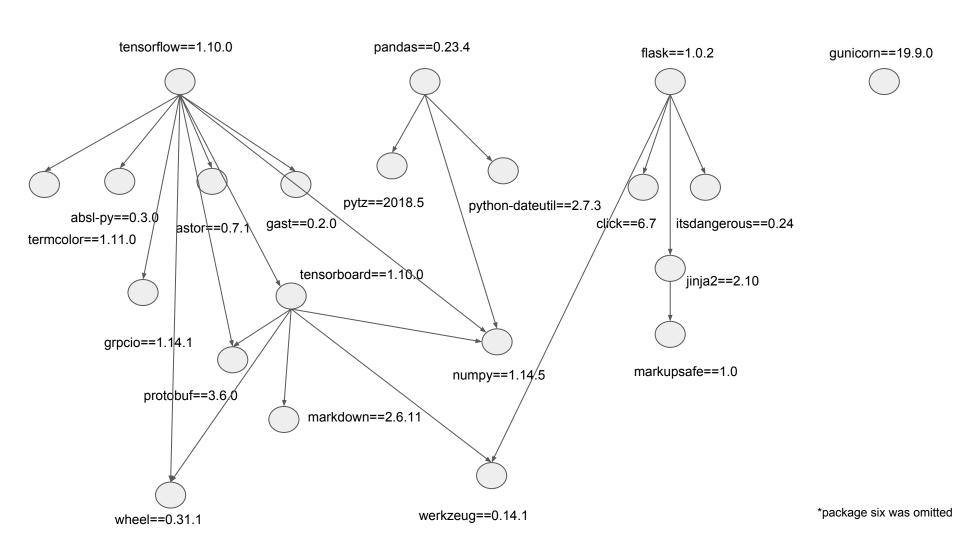


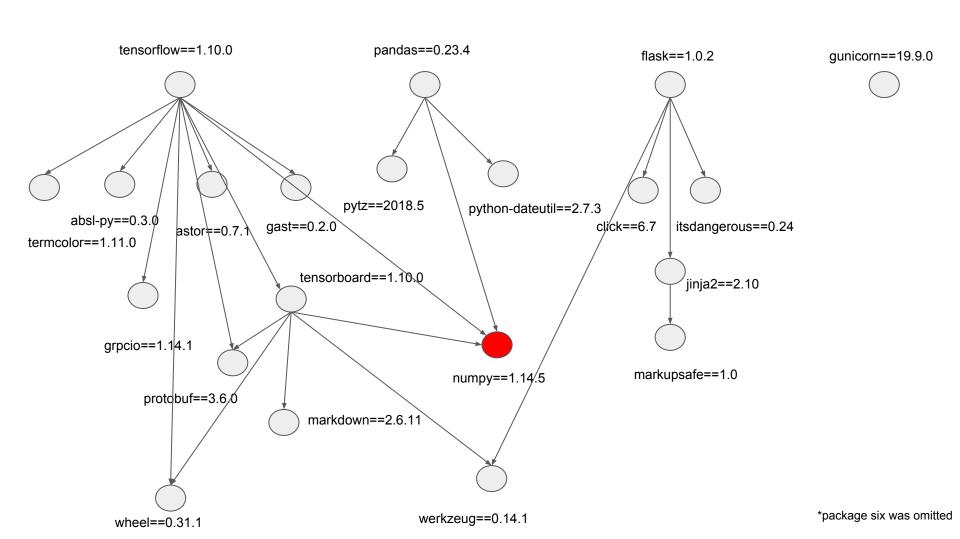


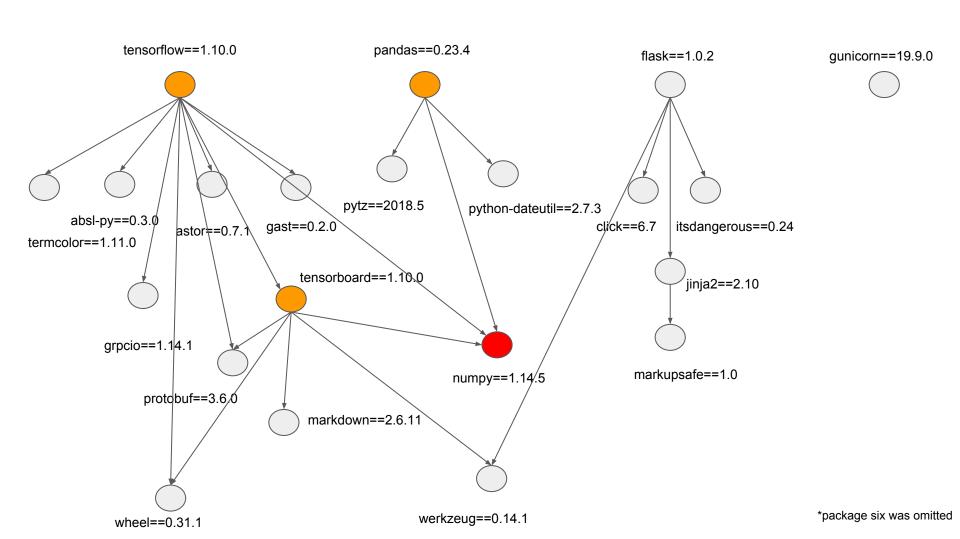


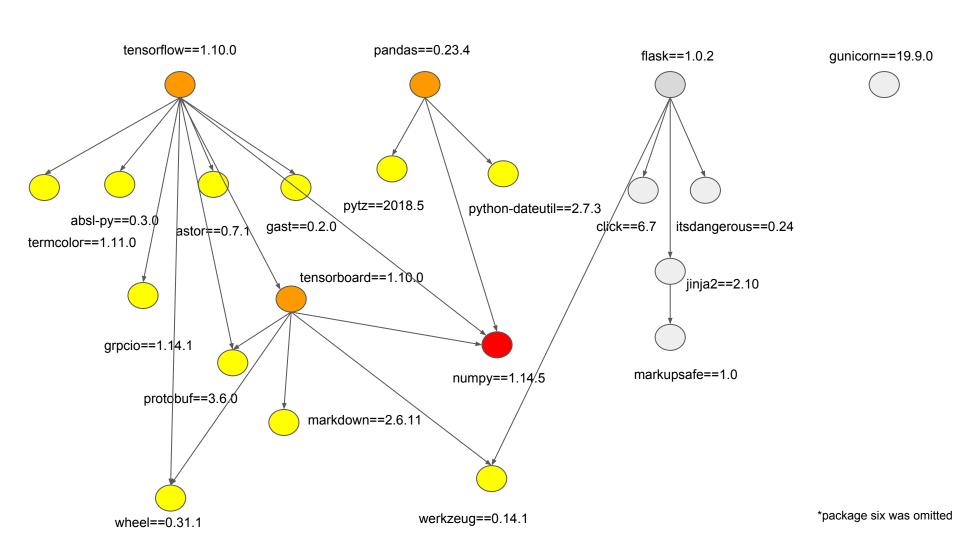


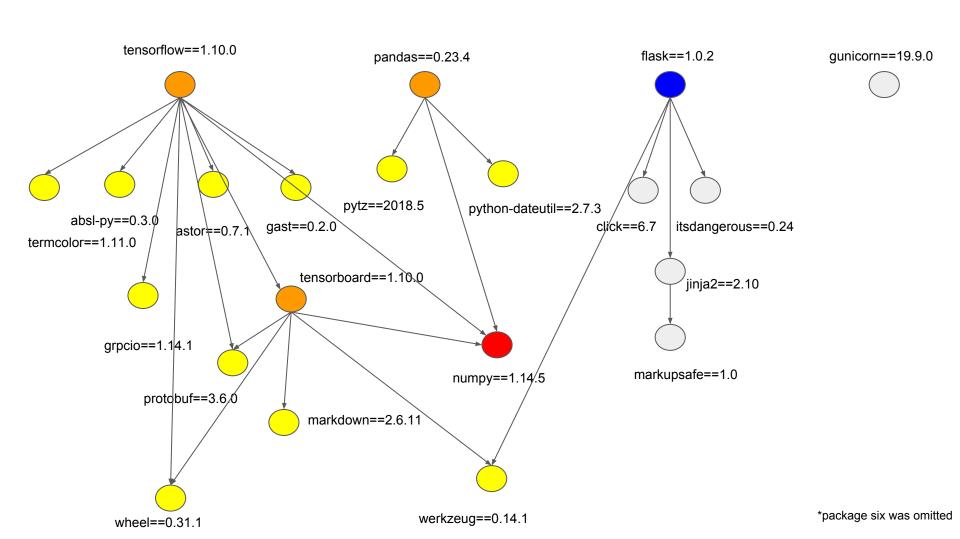


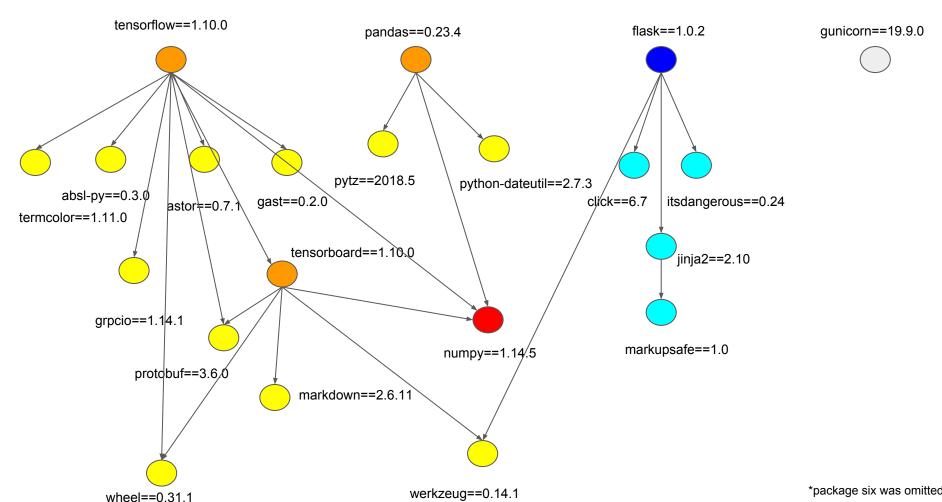












*package six was omitted

Thoth

- Server side resolution
- You state what you want, we what the resolved stack should look like
 - Based on experience we have
 - Based on observations we have
 - Based on monitoring we have

Thoth

- Server side resolution
- You state what you want, we what the resolved stack should look like
 - Based on experience we have
 - Based on observations we have
 - Based on monitoring we have

Knowledge

Your requirements for application

tensorflow >=1.5
pandas
flask
gunicorn

Thoth's recommendation on libraries you should use

- What is wrong
- What you should use
- Why you should use it

```
abs1-py==0.3.0
                           # tensorflow
astor==0.7.1
                           # tensorflow
click==6.7
                           # flask
flask==1.0.2
qast == 0.2.0
                           # tensorflow
grpcio==1.14.1
                           # tensorflow
qunicorn==19.9.0
itsdangerous==0.24
                           # flask
jinja2==2.10
                           # flask
markdown==2.6.11
                           # tensorboard
markupsafe==1.0
                           # jinja2
numpy = 1.14.5
                           # pandas, tensorboard, tensorflow
pandas==0.23.4
protobuf==3.6.0
                           # tensorboard, tensorflow
python-dateutil==2.7.3
                           # pandas
pytz==2018.5
                           # pandas
six = 1.11.0
                           # absl-py, grpcio, protobuf, python-dateutil, tensorboard,
tensorboard==1.10.0
                           # tensorflow
tensorflow==1.10.0
termcolor==1.1.0
                           # tensorflow
werkzeug==0.14.1
                           # flask, tensorboard
wheel==0.31.1
                           # tensorboard, tensorflow
```

Thoth

Your actual team member

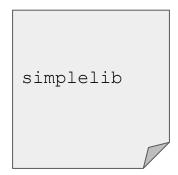
Bots proactively open pull requests

You review

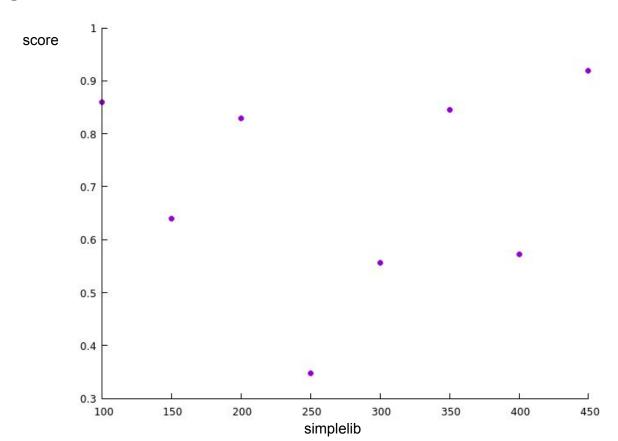
How to find the best software stack?

How good is my software stack?

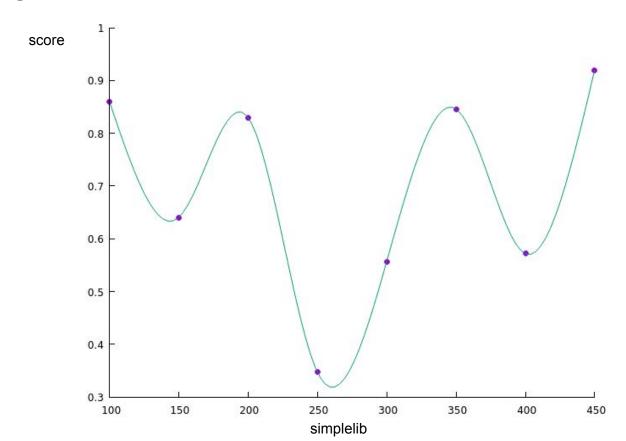
- Best = we need a mechanism to score software stack
- Let's have an application with one depndency
- This dependency has no dependencies
 - o no transitive/indirect dependencies



Scoring function - 2D

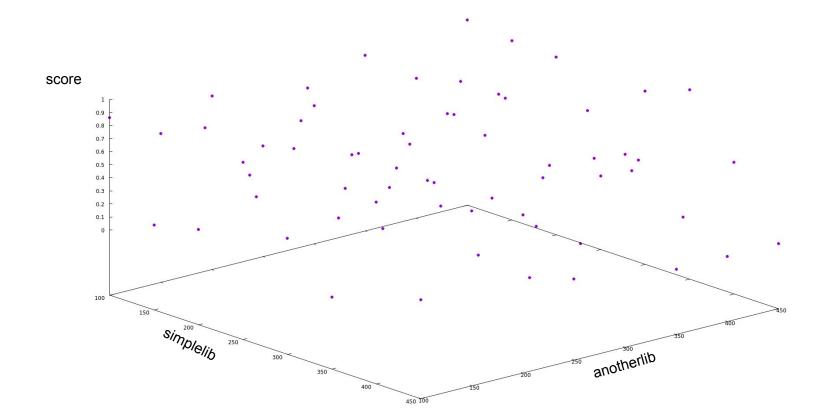


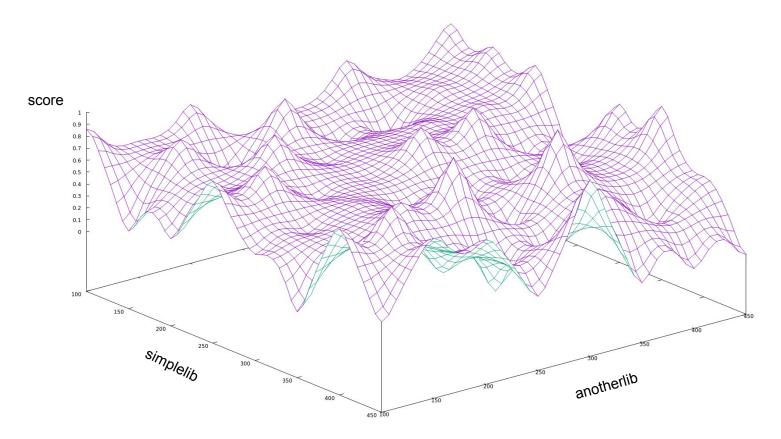
Scoring function - 2D

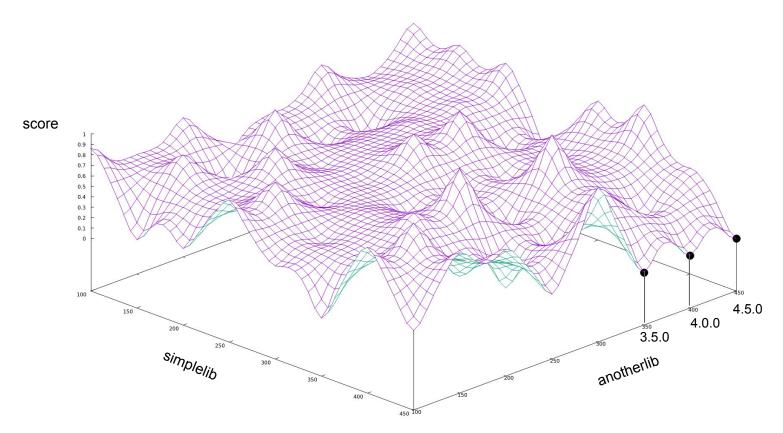


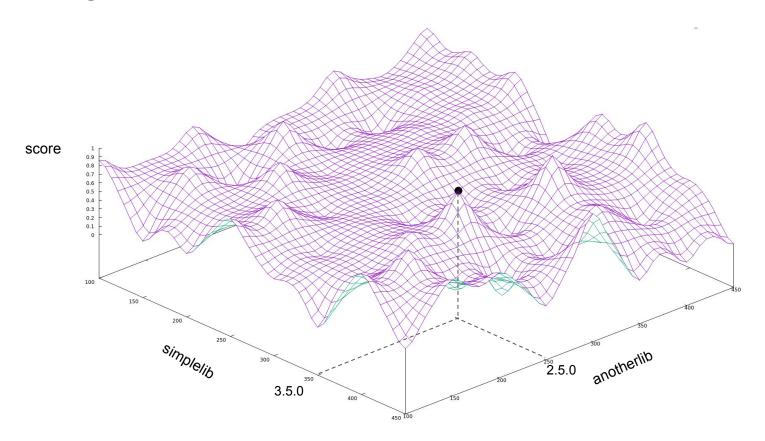
How good is my software stack?

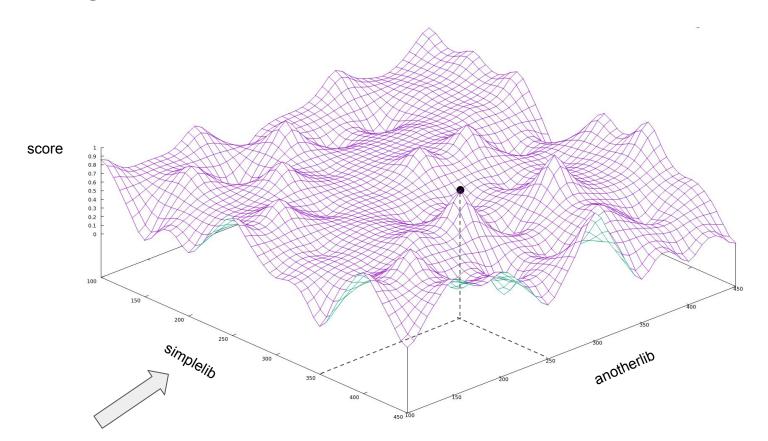
simplelib anotherlib

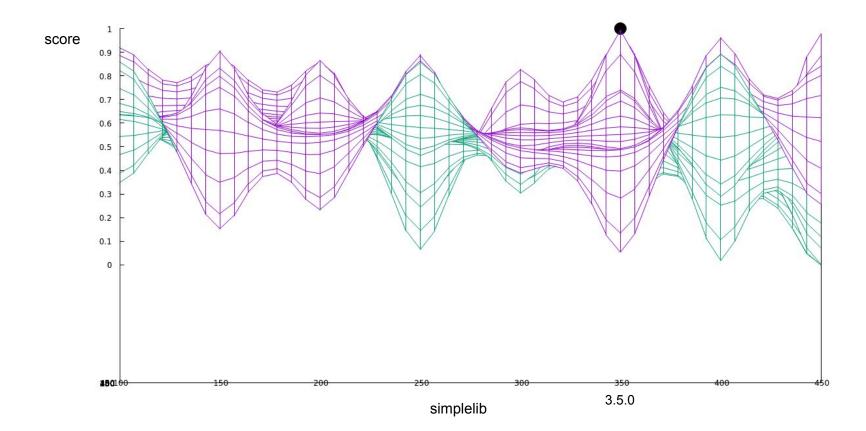












- n = number of all packages possible in my software stack + 1 (score)
- Scoring let's find a scoring function for my application
 - How does it perform when it comes to performance?
 - O How reliable are libraries?
 - How do libraries work together?
 - Do libraries have CVEs?
 - 0 ..

Scoring

- Package-level score
- Cross-library score

Knowledge

How to find the best software stack?

- Nearly impossible number of software stacks
 - Number of stacks with TensorFlow 68+ bilion possible fully pinned down software stacks
- However we can find "a good enough" software stack

Thoth's recommendation engine

- Resolve software stacks
- Score candidates based on knowledge
- Provide recommendation with reasoning
- Be close to source
 - Bots directly operate on source code
 - A user reviewes pull requests
 - If the pull request was closed => learn from it
 - Learn from opened pull requests and failures in CI
- Be proactive on changes around us
 - We monitor repositories
 - We proactively issue pull requests before something bad happens
- How to run your application

What does Thoth provide as of now?

- Automatic updates of your dependencies
 - Gathering information about software from logs
 - Does this update break your application
 - Log analyses
- Server side resolution
 - Finding package-level issues and excluding the given package from software stack
- Bots operating on source code
 - Automatic initial locks for your application
 - Automatic re-locking in case of issues
- Check provinence of your libraries
 - Am I really using well optimized TensorFlow on deployment?
 - Are packages installed from the configured source?
- Check your runtime environment

What is next?

- Aggregate more information about software stacks
 - Analyses of logs
- Add cross-library scoring
- Dependency Monkey
 - We already have dependency solver
 - Find out good TensorFlow configuration based on libraries and hardware used

Questions?

Thanks for your attention!

https://github.com/thoth-station