

NeurIPS - Thoth

Conference on Neural Information Processing Systems

Fridolin Pokorny <fridolin@redhat.com> & Thoth team

<https://thoth-station.ninja/>

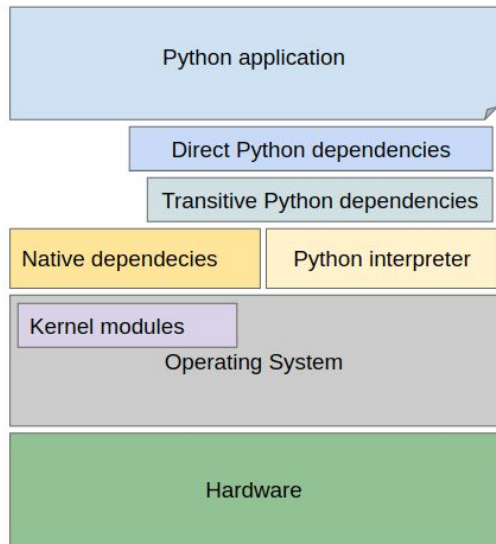
Red Hat AICoE, Project Thoth

2020




Thoth


- ▶ Why Thoth?
 - Large stack of dependencies to run an application
 - Each layer in the stack introduces complexity
 - Any component on any layer may cause issues
 - An issue on any layer creates an observable aspect
 - Performance
 - Security
 - ...
 - How to find the right stack for user needs?




Thoth: score of the stack




Thoth: score of the stack



Thoth: score of the stack



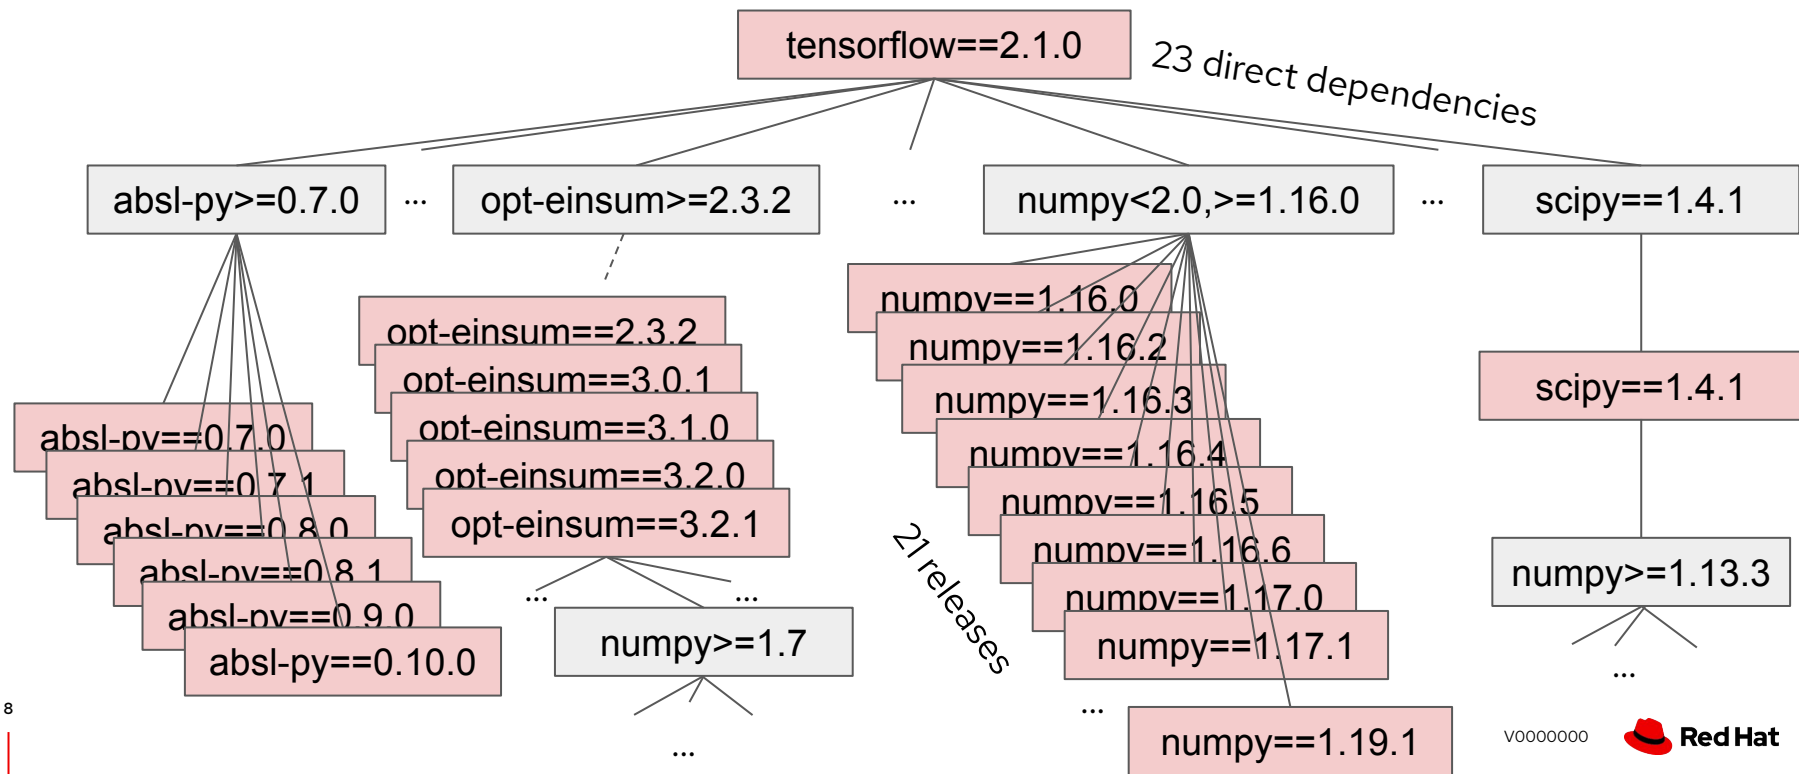
Thoth: score of the stack



Thoth


- ▶ Not just 2 dependencies but the whole dependency graph
 - Consider direct dependencies
 - Consider transitive dependencies
 - Consider runtime environment and other aspects of the software
- ▶ The state space is often too large to explore in real time
 - Number of combinations for a TensorFlow==2.1.0 stack: $\sim 3 \times 10^{13}$
 - Red Hat Developer: [AI software stack inspection with Thoth and TensorFlow](#)
- ▶ Can we split this into smaller sub-problems and find a solution to the optimization problem?

Thoth: TensorFlow dependency graph



Thoth: random state space sampling

- ▶ Randomly resolve software stacks based on the dependency graph
- ▶ Check “how good” they are based on knowledge about the software



Thoth: reinforcement learning

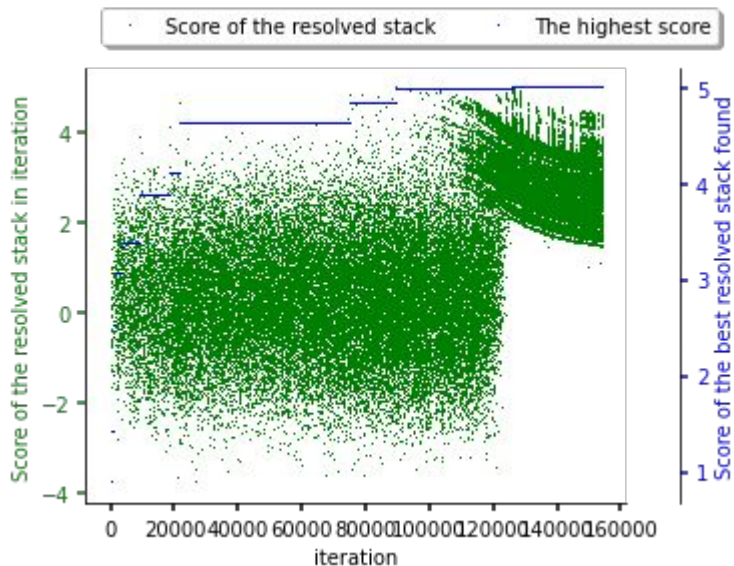

- ▶ Create a mechanism that can learn how to resolve software stacks
 - Using knowledge about the software to score software stacks
 - Learning how to navigate through the dependency graph to obtain best possible set of packages
- ▶ Gradient-free methods
 - Temporal Difference learning (TD-learning) and Monte Carlo Tree Search (MCTS)
 - Learning how to navigate through the state space and predicting score of trajectories in the dependency graph
 - Using “scoring pipeline” to score actions taken in the dependency graph and propagating a reward signal
 - Adaptive simulated annealing helps to balance exploration and exploitation

Thoth: example scenario

- ▶ Recommending a flask application
 - 7 direct packages considered in total
 - 7,861,340 possible valid software stack resolutions
 - Reinforcement learning based resolver run to resolve and score 70,000 software stacks (~1%)
 - Using randomly assigned scores for packages

Thoth: an example scenario

- ▶ Finding a software stack with a score of 4.99, the best possible candidate has 5.04




Call to action

- ▶ Use our data
 - Kaggle datasets at <https://www.kaggle.com/thothstation/datasets>
 - By conducting experiments on our datasets
 - By asking our services for advise
- ▶ Use our Services
 - GitHub App - <https://github.com/apps/qeb-hwt>
 - Thamos
 - OpenShift Pipeline Tasks
- ▶ Create Open Source software
 - Contribute a tiny feature (or a large)
 - [Project Planning is done openly](#)

Project Thoth

- ▶ AICoE, Office of the CTO
- ▶ Homepage
 - <http://thoth-station.ninja/>
- ▶ GitHub organization
 - <http://github.com/thoth-station/>
- ▶ Twitter account with updates
 - <https://twitter.com/thothstation>
- ▶ YouTube channel
 - https://www.youtube.com/channel/UCIUIDug_hQ6vlzmqM59B2Lw
 - <http://bit.ly/thoth-on-youtube>



Thanks for your attention!



<https://github.com/thoth-station>



<https://twitter.com/thothstation>



https://www.youtube.com/channel/UCIUIDuq_hQ6vlzmqM59B2Lw

References

Website <https://thoth-station.ninja/>

Twitter <https://twitter.com/thothstation>

GitHub <https://github.com/thoth-station>

 linkedin.com/company/red-hat

 youtube.com/user/RedHatVideos

 facebook.com/redhatinc

 twitter.com/RedHat