AI Center of Excellence
Office of the CTO

# Feedback about deployment of an intelligent app

DevConf US 2021

Francesco Murdaca
Senior Data Scientist



Red Hat

## Agenda

- Project Thoth
- MLOps Context
- Why, What, When
- How to use it
- Pipeline and results

# Project Thoth

Project Thoth

- Help developers in the selection of dependencies for their applications depending on their requirements
- Use bots to automate mundane work to offload humans work

# How Thoth can help developers?

- Keep dependencies up to date.
- Maintain software stack secure avoiding CVE.
- Recommend the most performant software stacks for AI Apps.
- Integrate source metadata information related to the packages used in the software stack to give advice to users.
- Integrate Thoth in day-to-day developers/data scientists tools.
- ...

# Thoth Integrations

- Command line tool thamos (developer laptop)
- Jupyter Tools (data scientist browser)
- Cyborg Kebechet (pull request/issues creator)
- Source-to-Image (container builder)
- Optimizing Deployment Pipeline

# MLOps context

# Providing feedback about deployment of an intelligent app

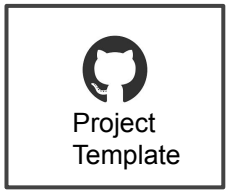



Prometheus

Grafana

## Providing feedback about deployment of an intelligent app

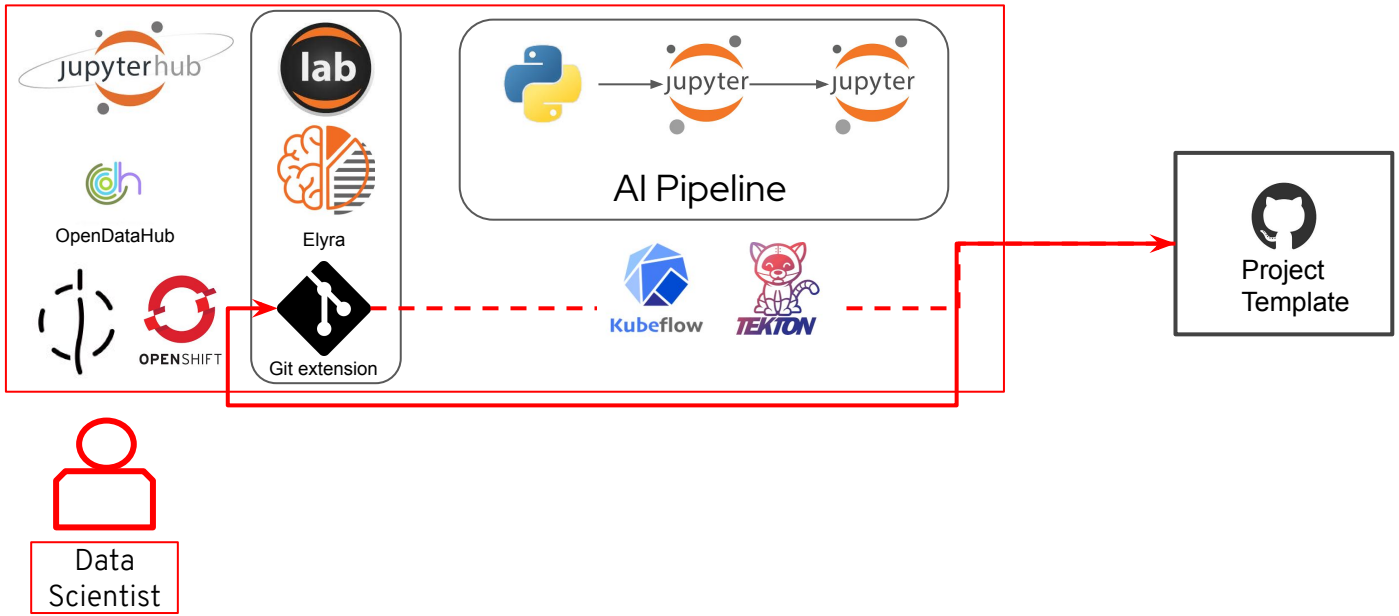




Prometheus
Grafana



Red Hat

# Providing feedback about deployment of an intelligent app

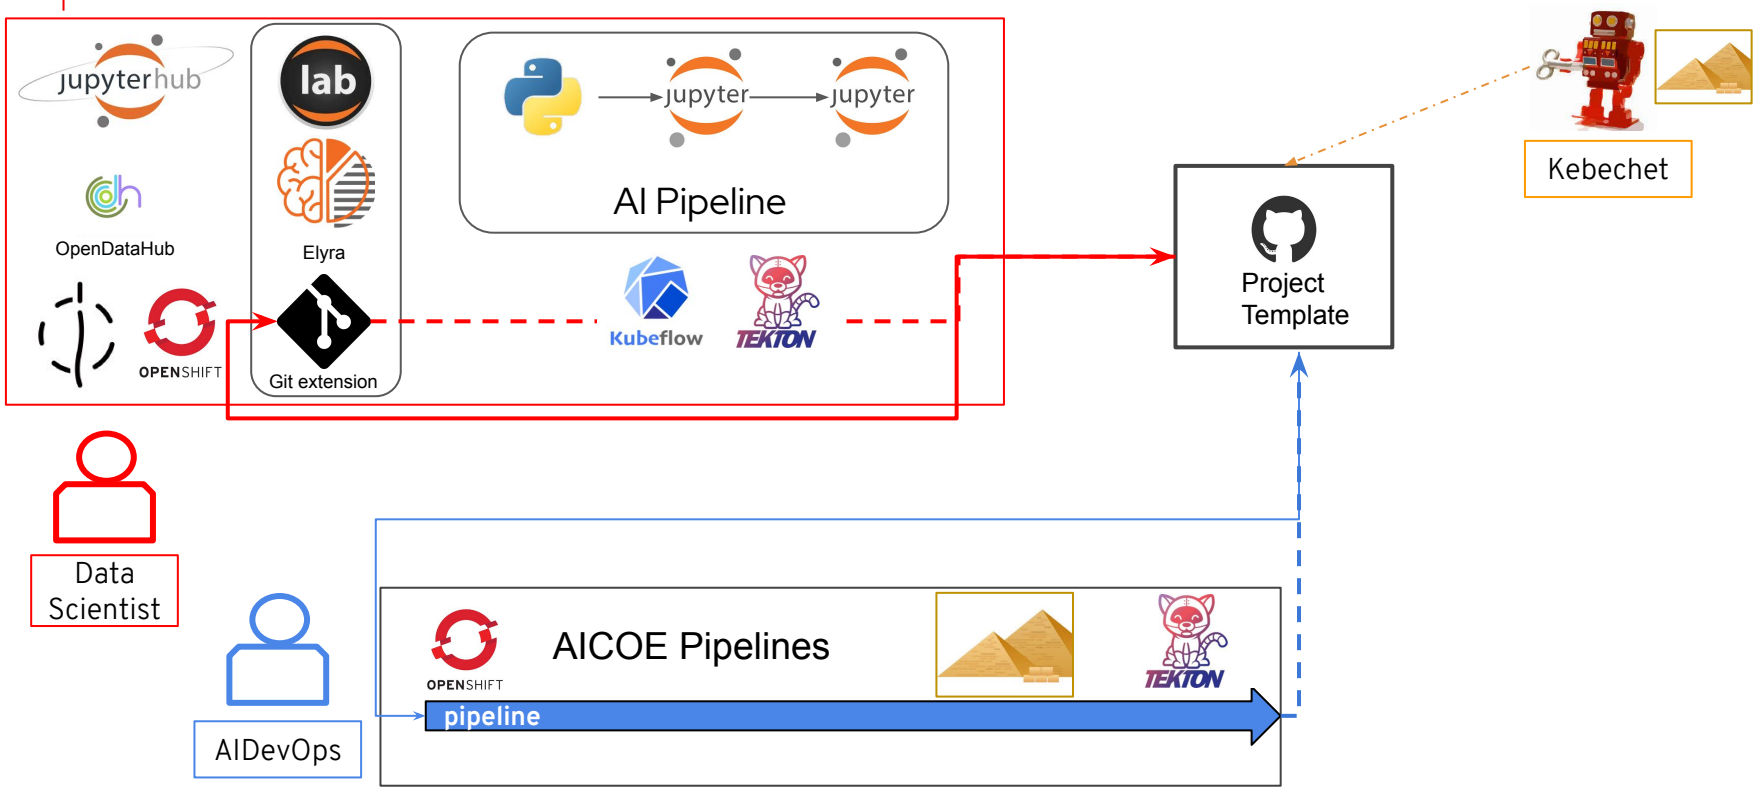






Red Hat

# Providing feedback about deployment of an intelligent app

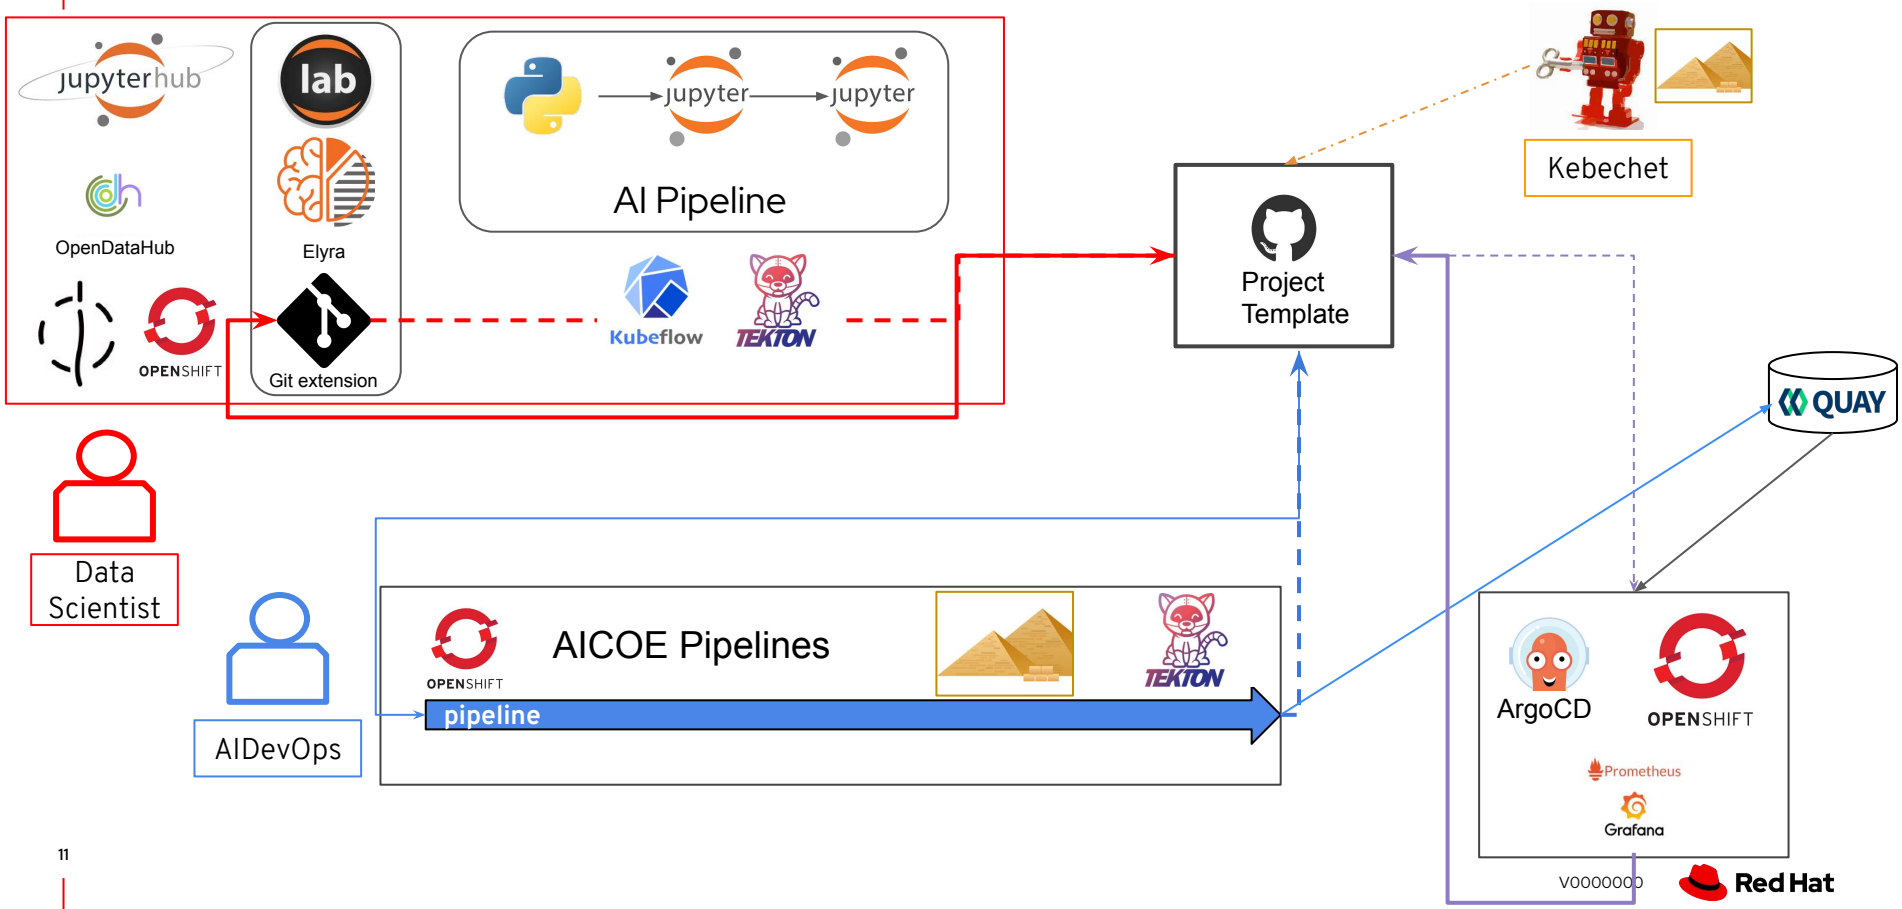

## Providing feedback about deployment of an intelligent app

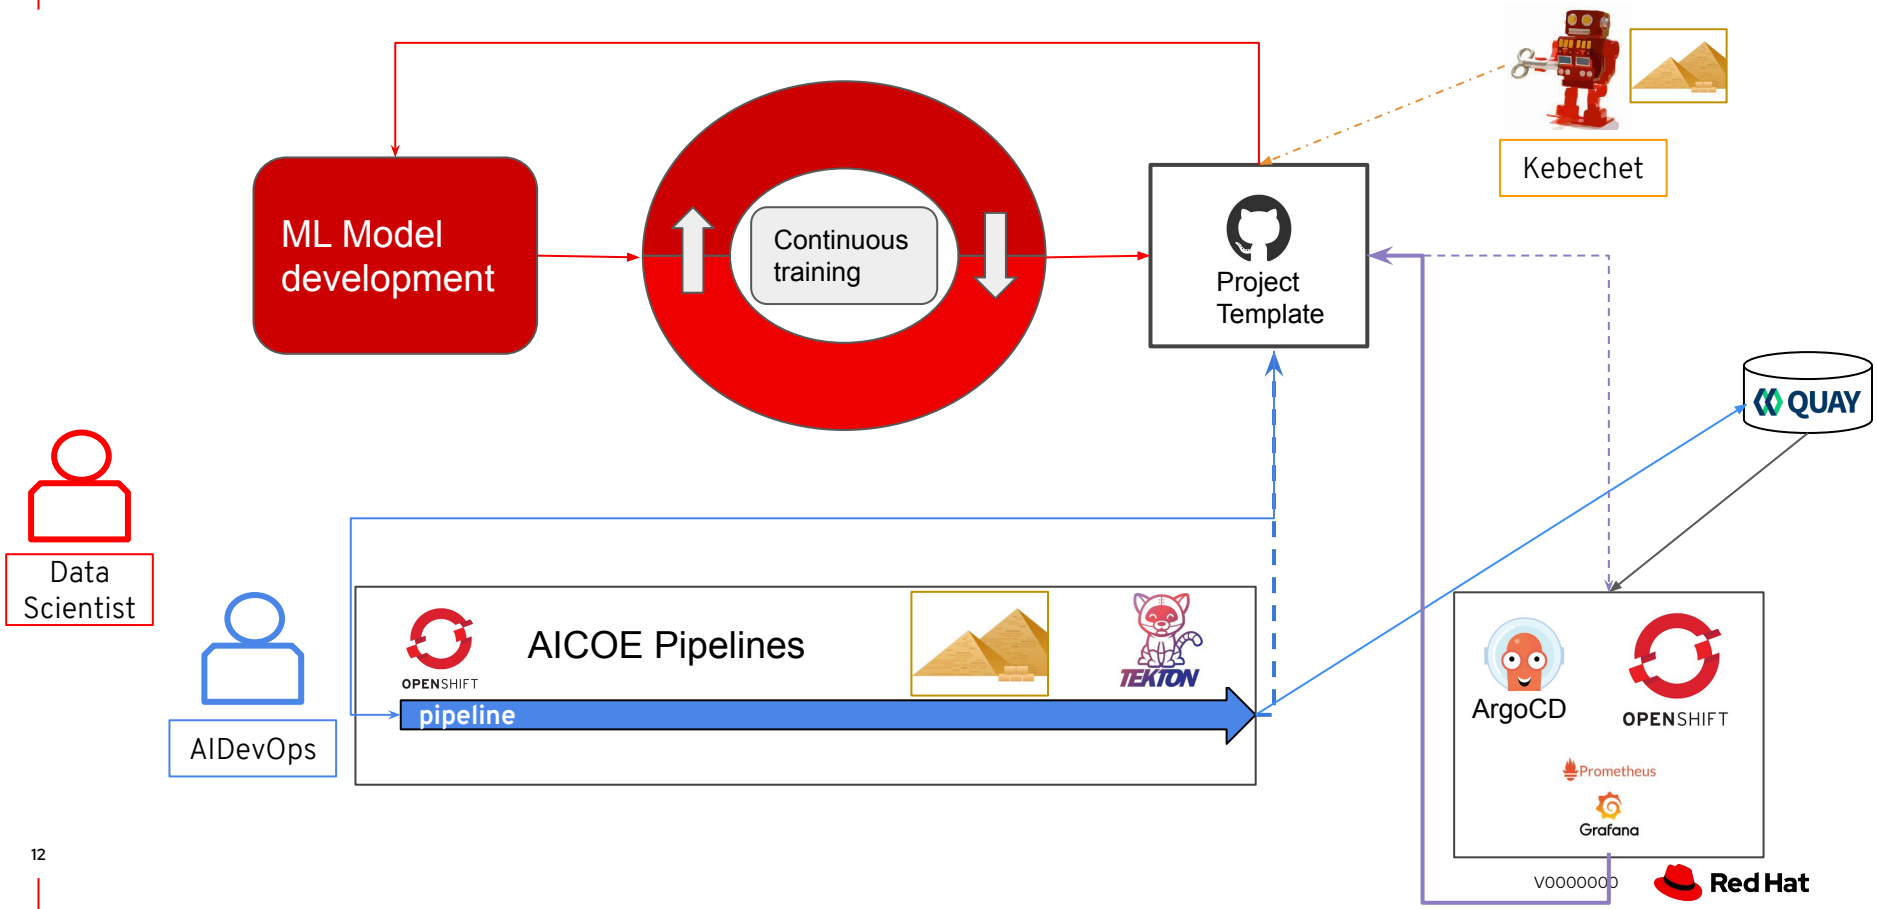

# Providing feedback about deployment of an intelligent app

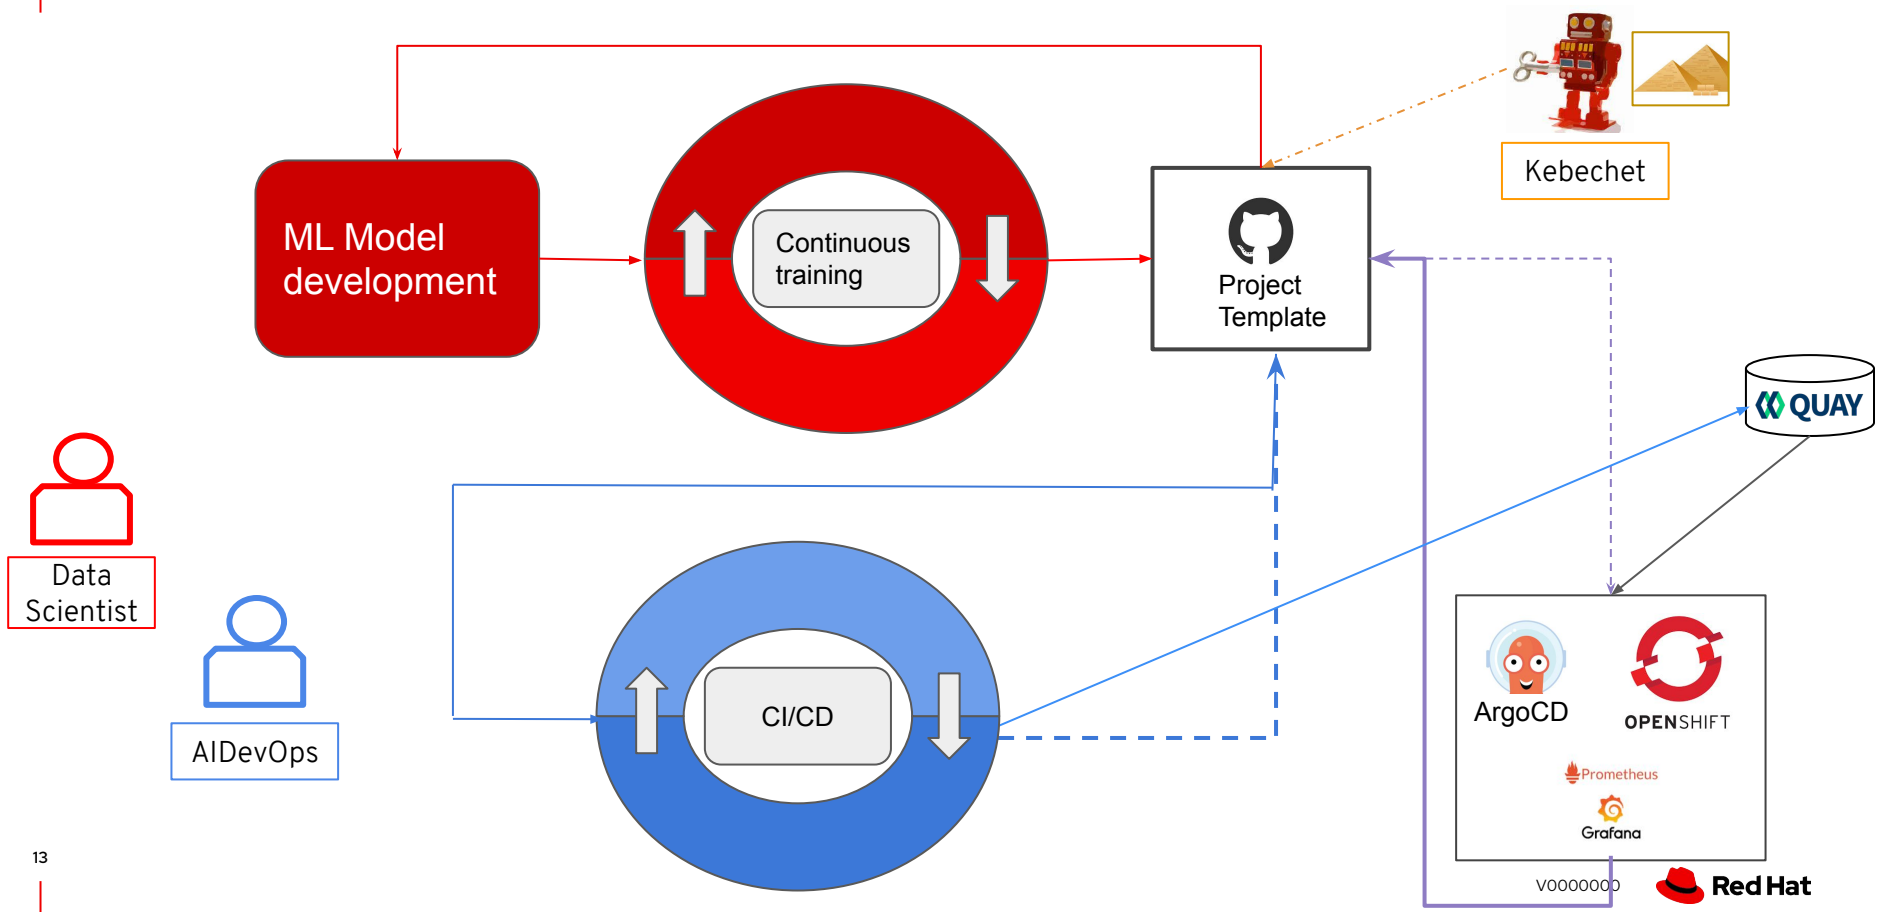

## Providing feedback about deployment of an intelligent app

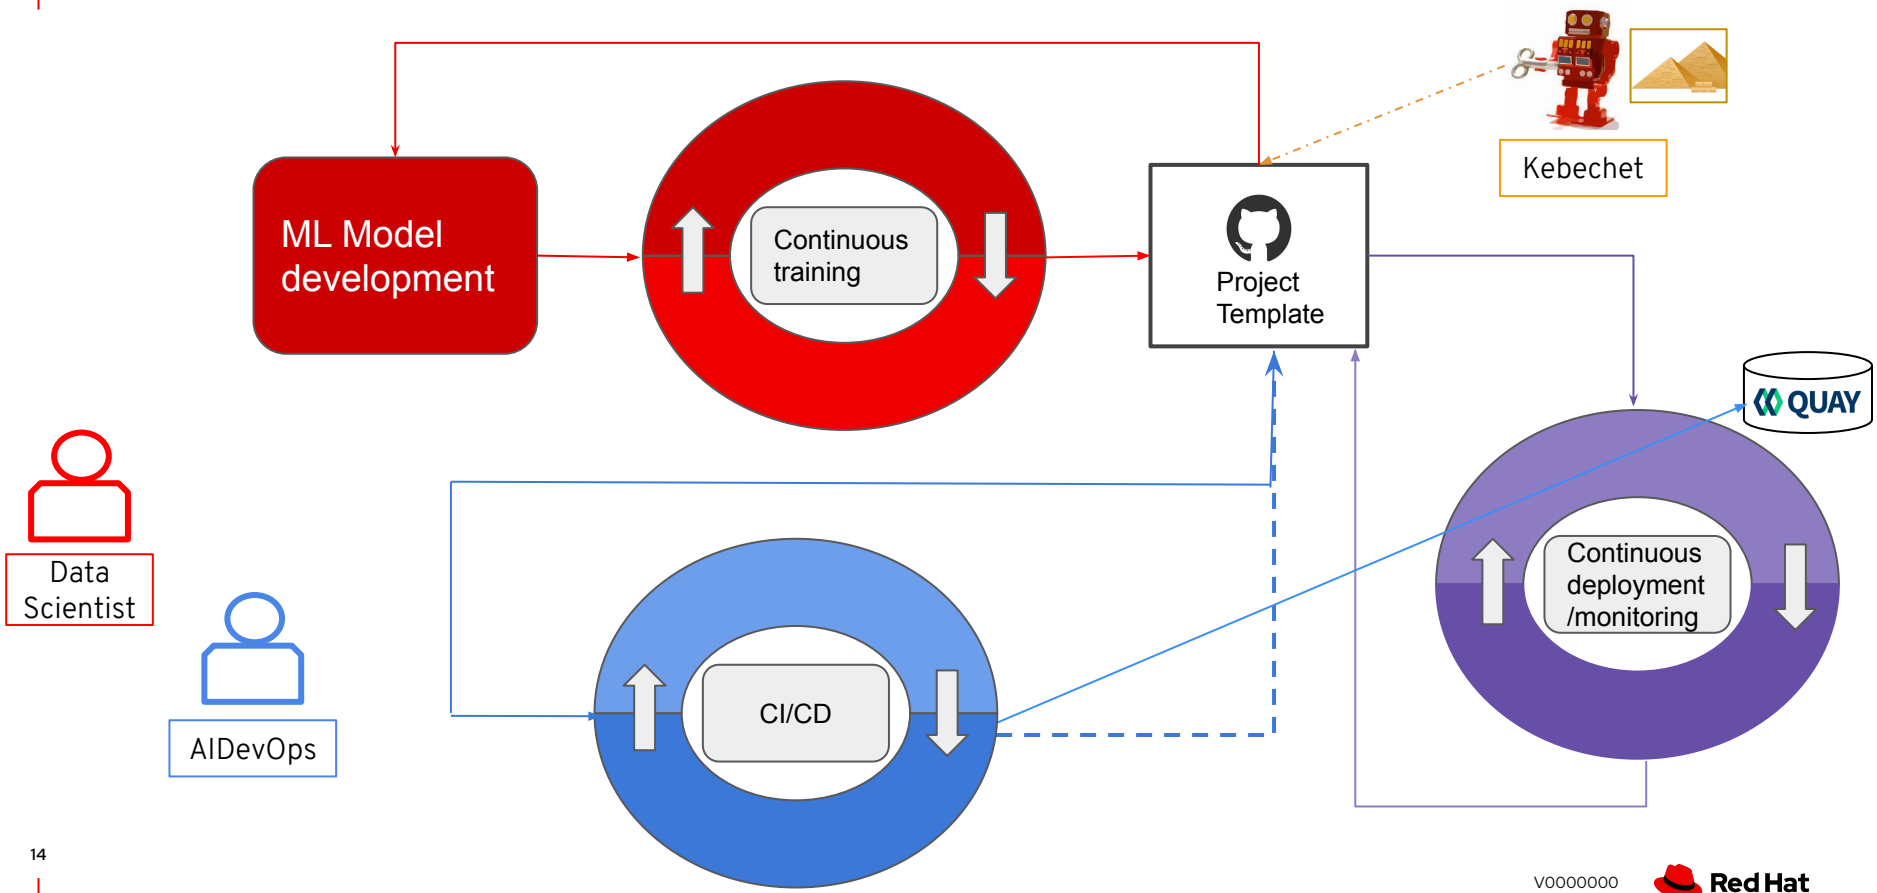






Red Hat

# Providing feedback about deployment of an intelligent app

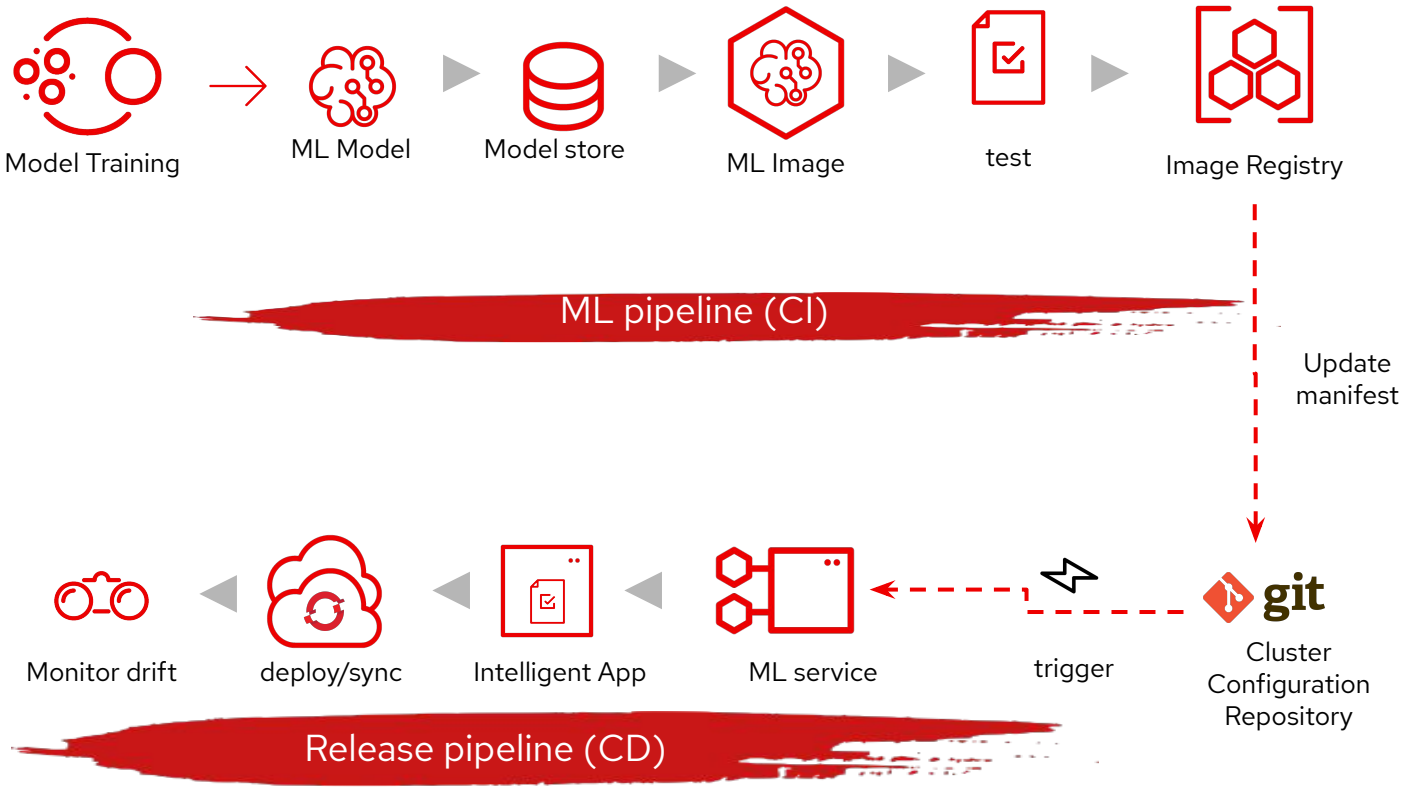

# Providing feedback about deployment of an intelligent app

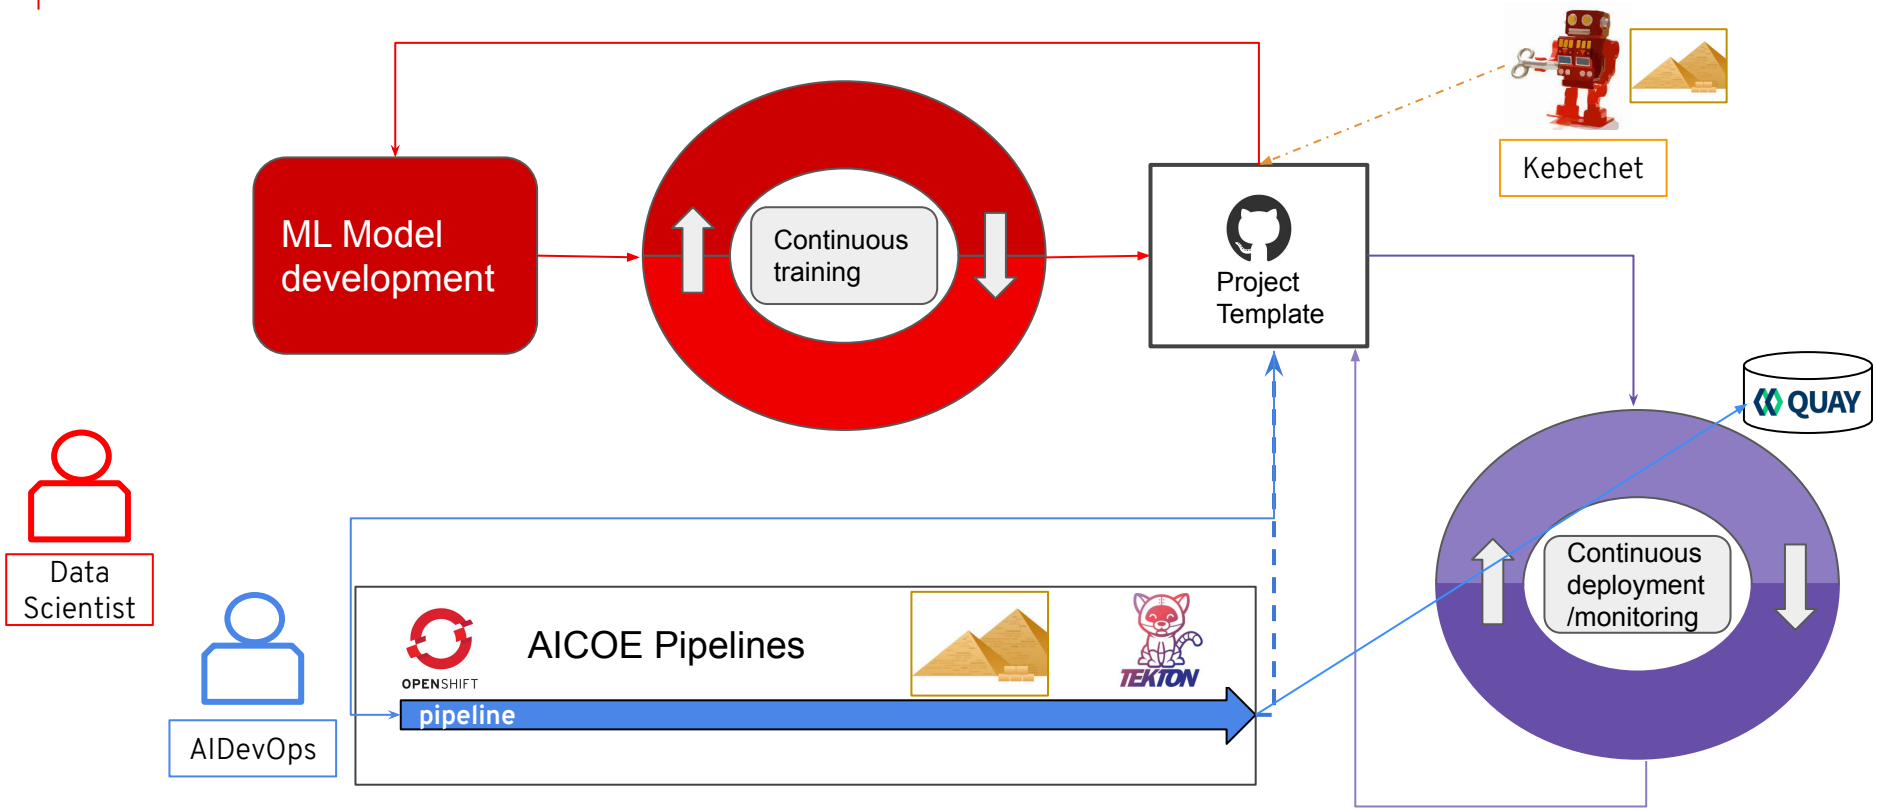






Red Hat

# Providing feedback about deployment of an intelligent app

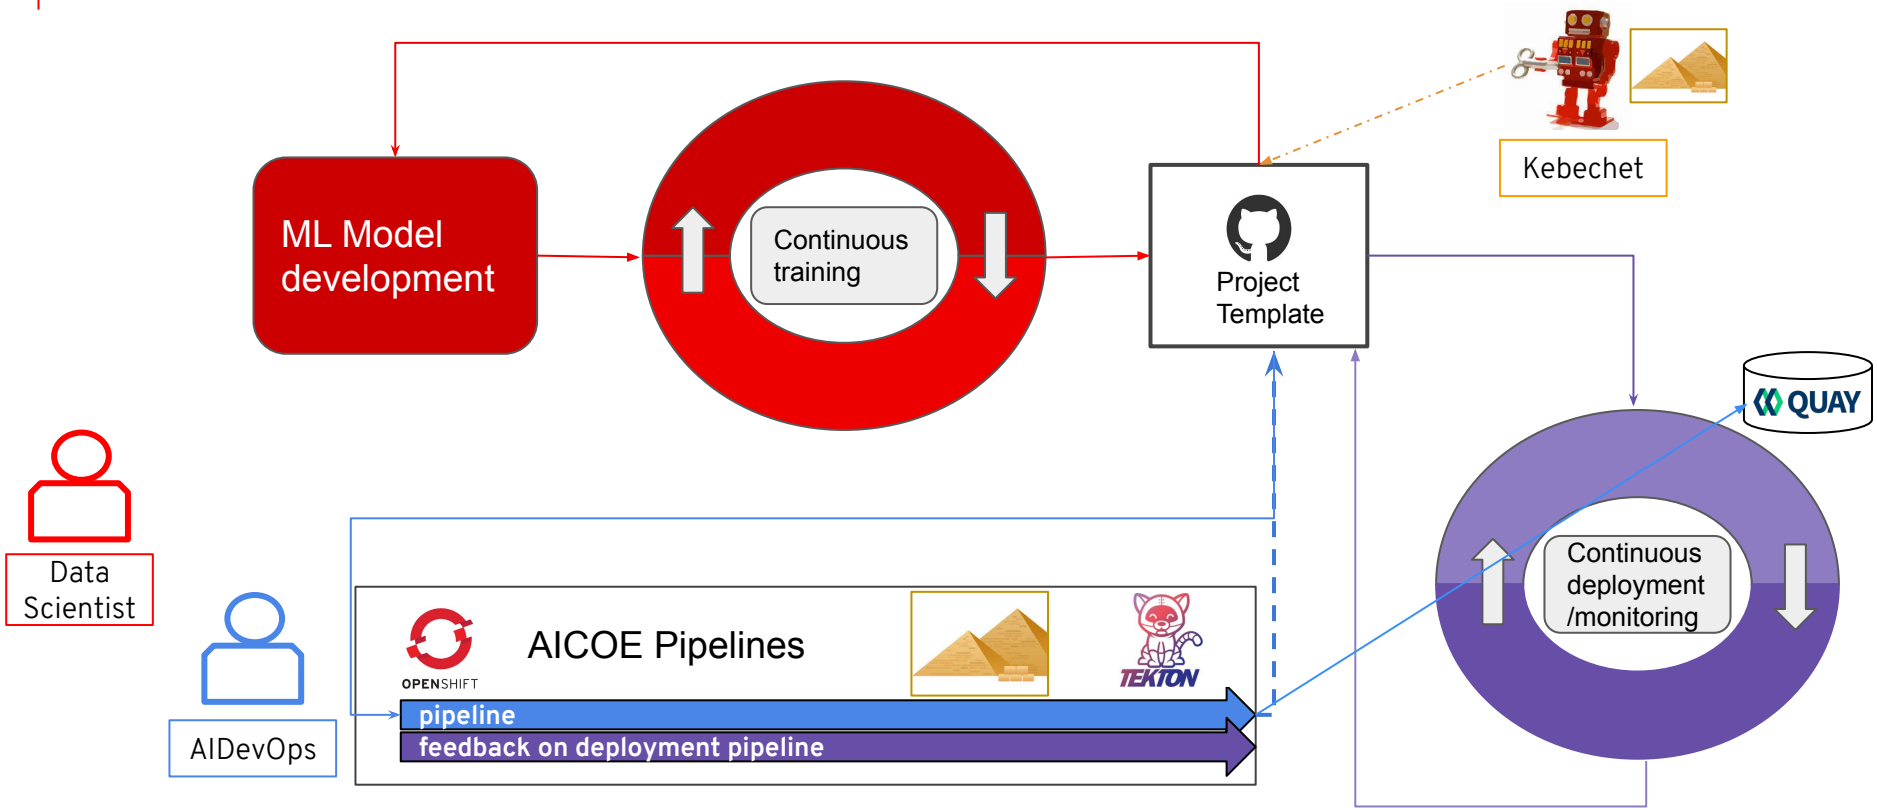






Red Hat

# Why, What, When

# Why feedback?

- App lifecycle is not static
- Speed up process
- Focus on modeling and data
- Reduce cost

# What feedback?

- Model metrics
- Application metrics
- Platform Metrics

# Personas that benefit from the feedback?

- AIDevOps wants to know operational information (latency, memory consumption, CPU usage)
- Data Scientist wants to verify model created is still performing well in production environment

# When giving feedback?

- When a new software stack is created
- When a new model version is created
- When a new dataset version is created

# How to use it?

# What do you need?

- AICoE-CI
- Model to be deployed
- Test to collect metrics

# AICoE CI

GitHub App

## aicoe-ci

This is the Continuous Integration Cyborg maintained by Thoth Station.

Configure

Manage your installation settings.

Developer

AICoE

Website

**aicoe-ci** is provided by a third-party and is governed by separate terms of service, privacy policy, and support documentation.

Report abuse

# AICoE CI

```
34    - name: inference
35      build:
36        base-image: "quay.io/thoth-station/s2i-thoth-ubi8-py38:v0.26.0"
37        build-stratergy: Source
38        registry: quay.io
39        registry-org: thoth-station
40        registry-project: elyra-aidevsecops-tutorial
41        registry-secret: thoth-station-thoth-pusher-secret
42      deploy:
43        project-org: "thoth-station"
44        project-name: "elyra-aidevsecops-tutorial"
45        image-name: "elyra-aidevsecops-tutorial"
46        overlay-contextpath: "manifests/overlays/test/imagestreamtag.yaml"
```



Source: https://github.com/AICoE/elyra-aidevsecops-tutorial/blob/master/docs/source/model-deployment/flask-application.md

# Deployments

- **Flask Application**
- KFServing
- Seldon
- TensoRT

# behave

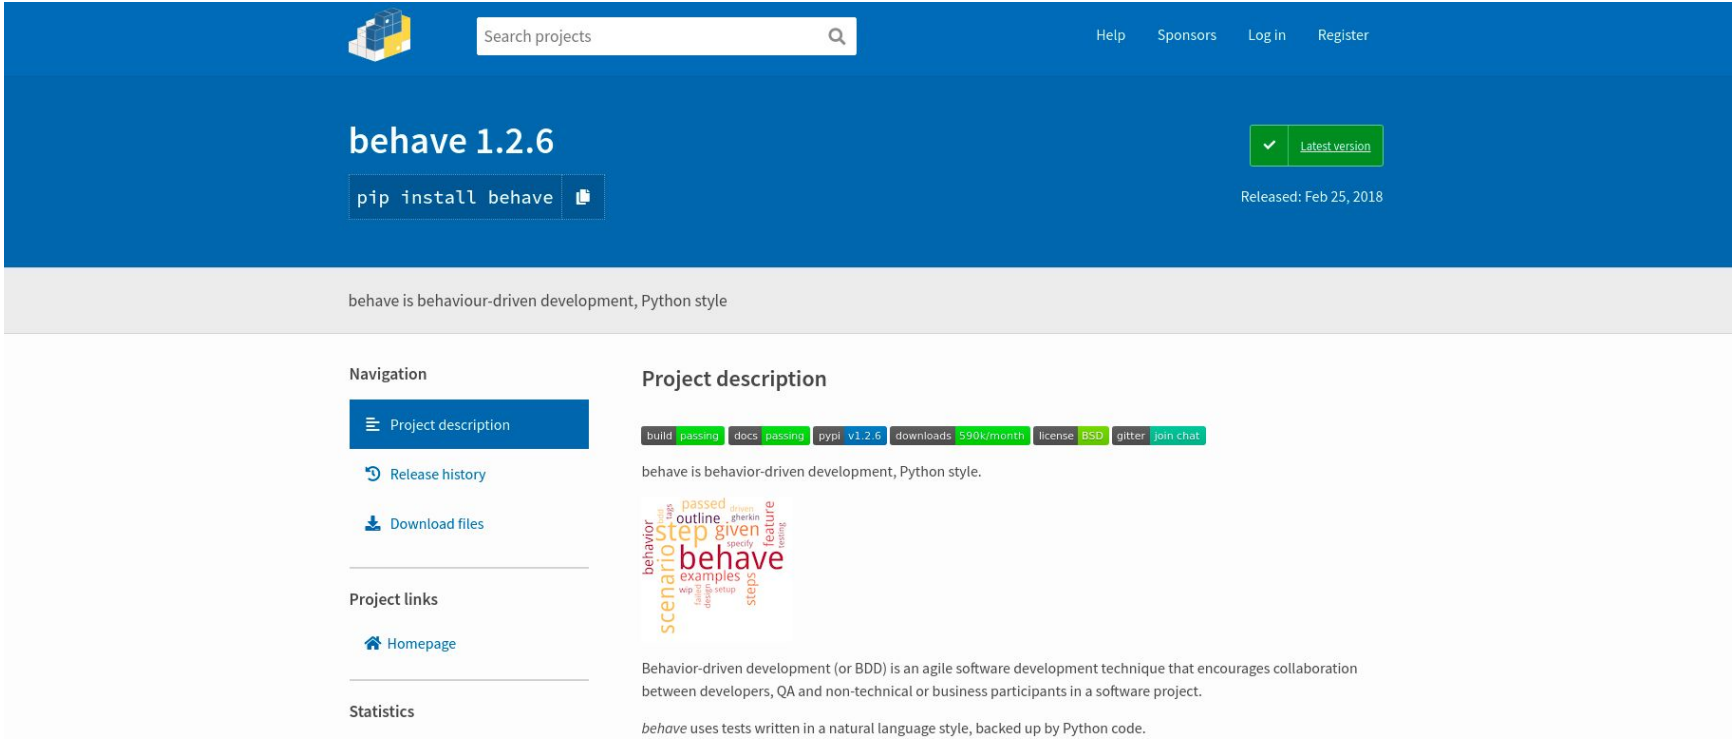




*behave* uses tests written in a natural language style, backed up by Python code.

## Providing feedback about deployment of an intelligent app

6 lines (6 sloc) | 274 Bytes

```
1  Feature: Gather metrics deployed model
2      Scenario: Deployment metrics gathering
3          Given dataset is available
4          Given deployment is accessible
5          When I run test to gather metrics on predict endpoint
6          Then I should get model and application metrics
```



Source: https://github.com/AICoE/elyra-aidevsecops-tutorial/blob/master/docs/source/test-model.md

# Providing feedback about deployment of an intelligent app

```python
@given("dataset is available")
def dataset_availability(context):
    """Check availability of dataset and retrieves it."""
    # Prepare MNIST data.
    _, (x_test, y_test) = tf_dataset.load_data()

    context.dataset = {
        "x_test": x_test,
        "y_test": y_test,
    }

    assert context.dataset


@given("deployment is accessible")
def deployment_accessible(context):
    """Check the deployment is accessible."""
    context.result = {}

    context.model_api_url = os.environ["DEPLOYED_MODEL_URL"]

    response = requests.get(f"{context.model_api_url}")

    assert (
        response.status_code == 200
    ), f"Invalid response when accessing {context.model_api_url}: {response.status_code!r}: {response.text}"

    assert response.text, f"Empty response from server for {context.model_api_url}"
```



Source: https://github.com/AICoE/elyra-aidevsecops-tutorial/blob/master/docs/source/test-model.md

# Pipeline and results

# Providing feedback about deployment of an intelligent app

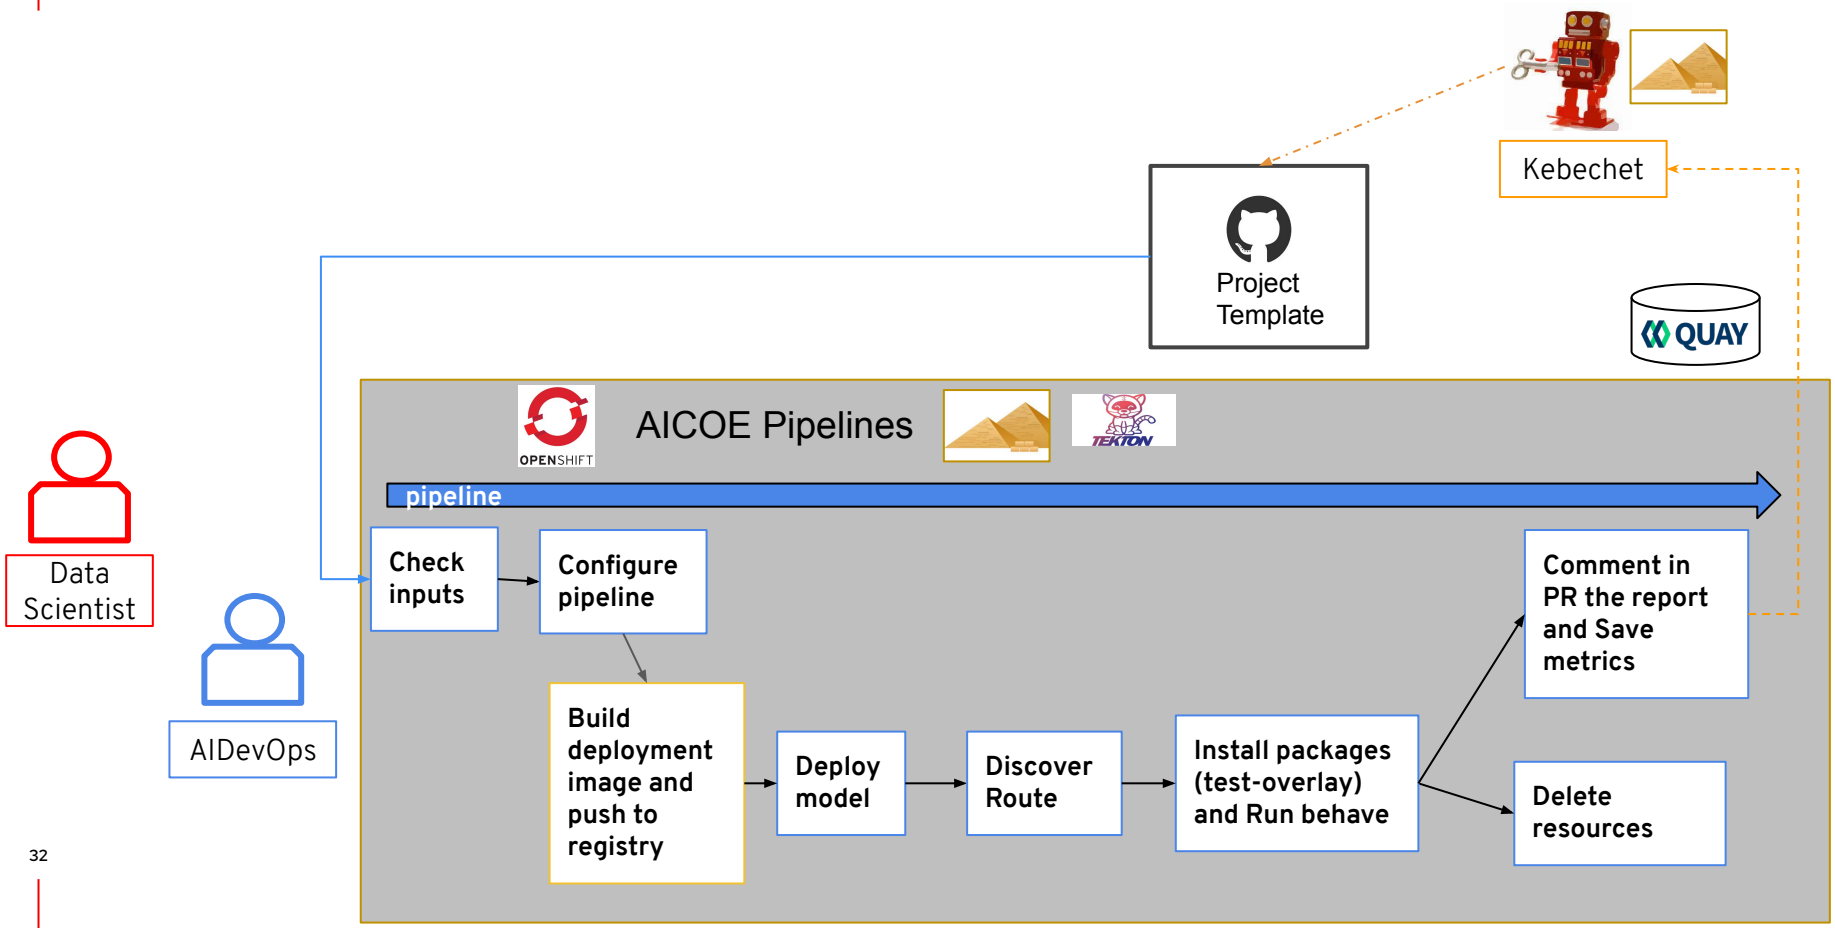

# Providing feedback about deployment of an intelligent app

**sesheta** commented 1 hour ago — Member

## AICoE CI results

### Test inputs

The following table shows info about test used to gather metrics.

| test URL | namespace deployment |
|---|---|
| https://github.com/AICoE/elyra-aidevsecops-tutorial/blob/e882fa5d4534aa0ec1909911665d1b05dab35736/features | aicoe-ci |

### Model and application metrics

The following table shows gathered metrics for model and application on your deployed models.

| average_latency | average_error | average_probability | number_requests | model_version |
|---|---|---|---|---|
| 0.118086 | 0 | 0.999916 | 21 | pr-160 |

### Platform metrics

The following table shows gathered metrics from platform on your deployed models.

| CPU max usage | Memory max usage |
|---|---|
| 1 | 384Mi |

- **Thoth station**
  - https://thoth-station.ninja/

- **AICoE CI**
  - https://github.com/aicoe/aicoe-ci

- **Thoth YouTube**
  - https://www.youtube.com/channel/UCIUIDuq_hQ6vlzmqM59B2Lw

# Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make
Red Hat a trusted adviser to the Fortune 500.


linkedin.com/company/red-hat


youtube.com/user/RedHatVideos


facebook.com/redhatinc


twitter.com/RedHat



Red Hat