

Thoth: healing Python applications

September 3rd, DevConf.US 2021

Fridolin Pokorny <fridolin@redhat.com>


\$ whoami && whoare

- Fridolin "*fridex*" Pokorny
 - Twitter: @fridex
 - Thoth team member since 2019
 - I like Python & road cycling
- Thoth
 - Started as a research project AlCoE team, Office of the CTO
 - One of the main offerings: cloud based Python resolver
 - <https://thoth-station.ninja>

The Python resolver run in cloud

- Recommendation engine for Python applications
- Publicly available to the community
- Stochastic resolver implementing gradient-free reinforcement learning methods
 - Temporal difference learning
- See documentation for more information:
 - <https://thoth-station.ninja/docs/developers/adviser>

Python cloud resolver



```
$ pip install thamos  
$ thamos config  
$ thamos advise
```

Why gradient-free reinforcement learning?

- Wide range (*infinite?*) of possible resolutions depending on requirements used in the application and other inputs to the resolver
- A model is trained on each request to the resolver
- Exploration phase and subsequent exploitation phase comes with the resolved software stack

Resolution pipeline

- The resolution process is using a pipeline made out of units of different type
 - Base pipeline types: boots, pseudonyms, sieves, steps, strides, wraps
- Pipeline units can be implemented directly in Python or declaratively in YAML files
- The resolution pipeline is constructed dynamically based on inputs to the resolution engine

Declarative interface for the resolver to resolve Python packages following prescribed rules

Prescriptions - declarative interface to the cloud based resolver

- Provide a way to declaratively state how the resolution process should look like
- A set of YAML files that are automatically consumed by resolver in a deployment
- See documentation for more information:
 - <https://thoth-station.ninja/docs/developers/adviser/prescription.html>

Prescriptions - Example 1

- Pillow in version 8.3.0 does not work with NumPy

<https://github.com/python-pillow/Pillow/issues/5571>

```
with PIL.Image.open(filepath) as img:  
    numpy.array(img, dtype=numpy.float32)
```

```
> frame_paletted = np.array(im, np.uint8)  
E   TypeError: __array__() takes 1 positional argument but 2 were given
```

```
/lib/python3.9/site-packages/imageio/plugins/pillow.py:745: TypeError
```

```
units:
steps:
- name: Pillow830TypeErrorStep
  type: step
  should_include:
    adviser_pipeline: true
  match:
    package_version:
      name: pillow
      version: ==8.3.0
      index_url: https://pypi.org/simple
    state:
      resolved_dependencies:
        - name: numpy
run:
  not_acceptable: Pillow in version 8.3.0 does not work with NumPy
  stack_info:
    - type: WARNING
      message: Pillow in version 8.3.0 does not work with NumPy
      link: https://github.com/python-pillow/Pillow/issues/5571
```

units:

steps:

- name: Pillow830TypeErrorStep

type: step

should_include:

adviser_pipeline: true

match:

package_version:

name: pillow

version: ==8.3.0

index_url: <https://pypi.org/simple>

state:

resolved_dependencies:

- name: numpy

run:

not_acceptable: Pillow in version 8.3.0 does not work with NumPy

stack_info:

- type: WARNING

message: Pillow in version 8.3.0 does not work with NumPy

link: <https://github.com/python-pillow/Pillow/issues/5571>

units:

steps:

- name: Pillow830TypeErrorStep
type: step

should_include:

adviser_pipeline: true

match:

package_version:

name: pillow

version: ==8.3.0

index_url: https://pypi.org/simple

state:

resolved_dependencies:

- name: numpy

run:

not_acceptable: Pillow in version 8.3.0 does not work with NumPy

stack_info:

- type: WARNING

message: Pillow in version 8.3.0 does not work with NumPy

link: <https://github.com/python-pillow/Pillow/issues/5571>

units:

steps:

- name: Pillow830TypeErrorStep

type: step

should_include:

adviser_pipeline: true

match:

package_version:

name: pillow

version: ==8.3.0

index_url: https://pypi.org/simple

state:

resolved_dependencies:

- name: numpy

run:

not_acceptable: Pillow in version 8.3.0 does not work with NumPy

stack_info:

- type: WARNING

message: Pillow in version 8.3.0 does not work with NumPy

link: <https://github.com/python-pillow/Pillow/issues/5571>

units:

steps:

- name: Pillow830TypeErrorStep

type: step

should_include:

adviser_pipeline: true

match:

package_version:

name: pillow

version: ==8.3.0

index_url: https://pypi.org/simple

state:

resolved_dependencies:

- name: numpy

run:

not_acceptable: Pillow in version 8.3.0 does not work with NumPy

stack_info:

- type: WARNING

message: Pillow in version 8.3.0 does not work with NumPy

link: <https://github.com/python-pillow/Pillow/issues/5571>

units:
steps:
- name: Pillow830TypeErrorStep
type: step
should_include:
 adviser_pipeline: true
match:
 package_version:
 name: pillow
 version: ==8.3.0
 index_url: https://pypi.org/simple
state:
 resolved_dependencies:
 - name: numpy

run:
 not_acceptable: Pillow in version 8.3.0 does not work with NumPy
 stack_info:
 - type: WARNING
 message: Pillow in version 8.3.0 does not work with NumPy
 link: <https://github.com/python-pillow/Pillow/issues/5571>

Prescriptions - Example 2

Adjust requirements in GPU enabled environments.

- Use tensorflow-gpu as a “*pseudonym*” to tensorflow if GPU enabled environment is available
 - [tf_gpu.yaml](#)
- Use the right tensorflow-gpu for the environment following support matrix
 - [tf_cuda.yaml](#)
 - [tf_cudnn.yaml](#)

Prescriptions - Example 3

Fixing library overpinning issues.

- TensorFlow in version 2.1 can cause runtime errors when running with h5py>=3 caused by overpinning
 - [tf_21_h5py.yaml](#)

Prescriptions - Example 4

Use a specific Python package index providing optimized wheel builds.

- Prioritize resolving AI CoE builds of TensorFlow for AVX2 enabled environments
 - [tf_avx2.yaml](#)
- Use *only* CUDA 11.1 builds of torch available on a PyTorch index:
 - [gpu_index.yaml](#)

Prescriptions - Example 5

Block using certain library functions due to security reasons.

- mktemp is deprecated due to vulnerability to race conditions
 - [tempfile.yaml](#)

Prescriptions - Example 6

Check configuration of the runtime environment. Pipeline units can also act on base image used.

- A GPU is available but no CUDA is available
 - [gpu_no_cuda.yaml](#)

Prescriptions - Example 7

Resolve considering RPM packages available in the runtime environment (similarly ABI or Python packages available).

- GitPython requires Git present in the runtime environment
 - [rpm.yaml](#)

References

- Prescriptions for Python open-source projects
 - Feel free to contribute to build better Python ecosystem:
 - <https://github.com/thoth-station/prescriptions>
- Declarative interface for the resolver to state requirements on Python packages in runtime environment
 - <https://thoth-station.ninja/docs/developers/adviser/prescription.html>
- Thoth homepage: <https://thoth-station.ninja>

Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



linkedin.com/company/red-hat



youtube.com/user/RedHatVideos



facebook.com/redhatinc



twitter.com/RedHat